

September 17, 2025

Oracle, Inference, and the Persistence of Disbelief

Jordi Visser

JVisser@22VResearch.com
22vresearch.com

Aside from Moonshots with Peter Diamandis, there are very few podcasts I listen to regularly. I'm much more of a thematic or guest-specific listener. In a world of too much information, podcasts are no different, so I need to be selective. Back in January, one theme dominated my listening: the DeepSeek moment.

When something impacts investor decisions and mindsets to the degree DeepSeek did, I want as much information as I can get on it. This particular theme ranged from Silicon Valley experts, traders, tech analysts, and pure quants. I wanted all their voices to understand if I was missing something. The problem wasn't finding opinions; it was filtering through the endless supply, as every single podcast was talking about it for weeks. Beyond the podcasts, the calls I received for my own AI opinion on the situation were greater than for any other event this year. It reached the point where I would guess that everyone in the investment industry remembers where they were when DeepSeek happened, even though there were no official numbers or factual information released. It was purely an embedded fear of the dreaded AI bubble.

Now fast forward and contrast that with last week, when Oracle dropped its latest earnings report. I'm not sure we've ever seen an earnings report like this. The numbers were staggering: Remaining Performance Obligations (RPO) surged to \$455 billion, up 359% year-over-year. Yes, I know, few of us knew what RPO was before last week, but now we do. Sometimes we big numbers don't shock us any more so to convert this into how large it is, if the New York Yankees had 50,000 seats and charged 500 dollars a seat and the 80 games were sold out every year since 1900 with those metrics, they would have \$252 billion in sales. For the last quarter alone, the Oracle sales increase was over \$300 billion. But the most important story for markets and GDP, in my view, was the breakdown of what is driving the growth.

Buried in that figure was the important macro point: inference demand has exploded. Training large language models may have dominated the AI narrative since the launch of ChatGPT, but it is inference, the actual deployment of AI into real workflows and the real economy, that is officially blowing apart expectations. I highlighted that we had reached this point in my May 15th report, [*The Inference Inflection: Where Real-Time AI Meets Real-World Opportunity*](#).

In that report, I listed 22 companies that should benefit from the move to inference from training, as enterprises begin using the thinking part of artificial intelligence to help make decisions. As of today since May 15th, an equal-weight basket of those names is up 26% vs. 13% for the SPX. The company earnings reports for Q1 showed that the inference demand was growing, but Oracle confirmed it had grown significantly. According to the earnings commentary for the company, one Oracle customer even requested all of the company's available inference capacity. That is not hype. That is not theory. That is real. As Larry Ellison put it on the call, it is "insatiable demand."

Inference demand should matter to every investor, strategist, analyst, and economist because it marks the seismic shift in AI from building models to using them, the point where AI moves from potential to real-world value creation. Remember this comes just weeks after an MIT report said AI had a "95% failure rate" that was intended to get clicks and scare people. Unlike training, which is episodic, inference is perpetual and scales with every interaction, making it the true driver of sustained adoption, productivity, and economic impact.

The sheer size of Oracle's order book puts this AI acceleration in perspective: the increase in backlog is nearly equivalent to the quarterly increase in U.S. nominal GDP. For a single company's contracted pipeline to rival macroeconomic output growth at the national level is extraordinary. And yet, despite this, the Oracle calls for my AI opinion never came like they did repeatedly for DeepSeek. In addition, I have yet to come across one podcast that dedicates more than a news

segment to its importance. The reality is the conversation about Oracle's numbers died much faster than when a Chinese company few people had heard of released a model to the public.

This is not surprising to someone who only speaks to investors about AI and where we usually spend most of the time with pushback. In this case, what makes this moment remarkable is not just Oracle's numbers, but the reaction to it. I thought for sure this would get everyone not on board to find a seat on this train, but even in the face of insanely large evidence, skepticism reigns. If you don't believe in AI, I understand why general economic, market, and even consumer sentiment won't improve. It is the economy and the stock market.

Last week, we got two examples of this negative sentiment. The AAI Investor Sentiment survey made the investor point clearly: last week's reading showed Bears at 49.5 compared to just 28 for Bulls. Across all of 2023 and 2024, there was only one week when Bears were higher than 49.5, November 2023, and that was after an 11% drop in the S&P 500. Also during last week, we got the University of Michigan Consumer Sentiment, and it came in at 55.4. In the last 45 years back to 1980, or a total of 540 releases, only 5 times have there been readings lower than the 55.4 reading. In those cases, the SPX had just fallen over 20%. This time it sits at all-time highs.

If I had to guess why so many people remain bearish, it is simply that most people do not believe in AI. Last week, the cover of The Economist had an AI title, and everyone acted like this was it, the curse of The Economist cover article, but the article according to ChatGPT was a 4 on a 1-10 Bearish to Bullish scale. Regardless of the numbers from Oracle or Nvidia or the Mag7's 26% YoY earnings growth, there is still a feeling that AI is just hype that will eventually fail. If you take away AI, there is no economy, so I actually agree that if you leave everything else the same, the economy and market would be in trouble. But that is not the case and not what is happening.

As I have my conversations around AI each week, I can say they almost never start with anything positive. The ones making money on it are still focused on what can go wrong. The ones not making money are argumentative about AI itself and what it is actually doing. I consider myself a curious person who loves to learn about new things and go deep into rabbit holes about the future. This is why I have always been a macro person. I want to understand in my mind what five years will look like and work backwards. I believe in five years we will be curing all diseases and exponentially expanding lifespan, have flying cars from NYC to JFK so we never have to be on the Van Wyck again, or that we will get on an elevator with a humanoid next to us, and so many other mind-bending things. The economy we know today will not exist. But for many, they don't see it that way.

This is my attempt to now take the Oracle numbers and be understanding as to why there is a blind spot. Macro investors have struggled to capture the real impact of AI on GDP, and I agree there is no real way to measure it the traditional way GDP was created. Simon Kuznets, the father of the GDP measure, did not know AI would be happening or that we would still be using it, so forgive him for not focusing on the intangibles like software, the cloud, and now intelligence embedded in machines. Joseph Schumpeter saw it with Creative Destruction. Part of the problem is that AI's effects are multi-layered and don't fit cleanly into traditional macroeconomic models. It especially breaks the relationship between labor and GDP. Here is my view on what has happened for its impact on the way GDP is measured.

First Wave: Infrastructure Buildout

The early impact showed up in the surge of capex for chips, especially GPUs, which powered the training of large training models. Also, the smaller construction of data centers that have occurred so far. That was the spark.

Second Wave: Margins and Earnings

AI's effect on corporate profitability has already begun. This began with the personal computer, the internet, and the cloud. By driving operating efficiencies and boosting top-line growth, those prior innovation advances and now AI have helped sustain record profit margins. That, in turn, has kept equities climbing and consumption resilient through the wealth effect and a top-heavy concentration of wealth and spending, even in the face of higher rates. Back to U of Mich Sentiment, in 2019 before Covid, we were near the highs of Sentiment with 116 trillion dollars in Household net worth and now five years later with net worth at 176 trillion dollars, sentiment is near all-time lows. AI and tech have caused this issue and kept the spending going.

Third Wave: Data Centers and Power

Now we are entering the phase where data center buildout and power infrastructure will scale to meet Oracle-sized orders. This is where AI's impact extends beyond tech into industrials, utilities, and commodities. This is why I think PMI is going to rise as a tide lifts all boats because Oracle inference demand outstrips the supply chain of compute and power.

Fourth Wave: Digital and Physical Employees

However, from an economic basis, the next five years will see the most profound shift: the introduction of digital employees, AI agents, and the early adoption of physical employees in the form of robotaxis and humanoids. This will supercharge productivity. Digital agents will automate knowledge work. Robotaxis will reshape transportation economics. Humanoids will begin to handle repetitive physical labor. The multiplier effect on GDP will be immense. This is where if the Fed is focused on labor, expect to see lower rates while GDP is good and stocks are rising like now.

It is important to remind people that this entire expansion into AI growth is being funded by the companies (Mag7) who creatively destroyed all the other companies without issuing debt. Technological innovation made the 2010-2019 was also a period where macroeconomists could not understand the economy which is why they changed its calculation significantly. In the 1990s after the birth of the internet people said to not fade the US consumer who kept spending and again, the GDP calculation was changed. AI is just a continuation, at a much faster and greater scale, than the disruption that occurred in the prior 30 years.

Economists trained to measure output through traditional factors, labor, capital, and productivity, as an output of the industrial focused built GDP index are missing the exponential layering happening here. GDP is a useless measurement in a world dominated by AI. Oracle's earnings are not just a corporate story; they are a macro data point, a flashing signal of structural economic change. I will bet that within the next decade, GDP will be thrown out as a measurement as the digital economy accelerates and calls for a new index grow due to the continued rise in inequality.

The natural extension of inference is the AI agent. These are not chatbots spitting out generic answers. Agents reason, plan, and execute multi-step tasks in real-time and around the clock. Their hourly earnings and hours worked are fixed. Imagine an agent combing through a company's private data, analyzing a client's purchase history, cross-checking it against market conditions, and producing a hyper-personalized sales strategy on the fly. That is not five years away; that is the path Oracle's customers are signing up for today.

Agents require constant streams of inference compute. They don't pause; they act. And as they act, they multiply the demand for capacity. This is why Oracle's backlog ballooned. It isn't just about serving queries. It is about building the infrastructure for a future where AI agents orchestrate entire workflows.

If Oracle is reporting these numbers, imagine what AWS, Microsoft Azure, and Google Cloud are about to show for their cloud revenues. Oracle's backlog growth validates what is happening across the entire market. Cloud revenue numbers in upcoming earnings should be strong, if not spectacular. Those stocks may have already discounted the growth, or their spending may create investor doubt in their ability to meet revenue needs, but I doubt their earnings will not be able to grow with these types of indications of demand.

And the multi-cloud strategy adds another layer. Oracle executives noted their database services embedded in Microsoft, Google, and Amazon are growing at "incredible" rates. Competition is giving way to collaboration, because the scale of demand is too big for any single provider to capture alone. This is not a zero-sum game; it is a rising tide.

The disbelief in markets today mirrors the disbelief that always accompanies technological revolutions. When electricity spread, when the internet arrived, when smartphones reshaped behavior, there were always reasons to dismiss, to doubt, to delay belief. AI is following the same trajectory, only faster. According to the Silicon Valley experts close to AI, there is a decade of innovation happening each year, and it compounds so it is accelerating as I write this paper. Oracle is an indication of that.

Oracle's earnings are the canary in the coal mine or, more accurately, the thunderclap that should wake people up. Inference is the bottleneck. Enterprise adoption is accelerating. AI agents are coming. And hyperscalers are about to print massive cloud growth. Behind that are autonomous vehicles, flying taxis, and humanoids.

I am not a Star Wars or Star Trek kid; I hate both fantasy books and movies. I think I am the only person who could not watch Game of Thrones due to dragons. Despite this, as a macro person who uses it every day in ways that multiply as each model is released, this is not fantasy. This is reality, and to fade it, in my opinion, is a mistake that will chew up every investor, employee, company, country and children of the parents who don't support them using it through its massive disruption.

Sentiment remains mired in denial. That denial will not last forever. The Oracle contracts are signed, the workloads are growing, and the demand is real. By the time disbelief finally gives way to recognition, the opportunity may already be priced in. For now, it is not. Along the way there will be more DeepSeek panics and doubts that lead to 20% corrections in the SPX. The inequality and government debt used to try and keep the inequality anger in check, will lead to episodic falls in stocks. But AI's destiny will not be denied as it is a global race.

January's DeepSeek panic was an indication of how many people doubt AI. Oracle's September dismissal is the second. The market can cling to disbelief, but reality is now visible in the numbers. AI is not a story about the future; it is a story about the present.

The irony is that disbelief itself has become the trade. The more the market resists the evidence, the greater the opportunity for those willing to accept it. Retail investors remaind willing to accept it. Oracle's quarter is not an isolated event. It is a signal of what is to come, a world where inference capacity, AI agents, and multi-cloud infrastructure redefine how businesses operate and how investors should position.

And this time, the orders are not just big. They are the size of GDP growth itself.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.