

October 7, 2025

From the Playbook to the Play: Why NPUs Are AI's Championship Moment

Jordi Visser

JVisser@22VResearch.com

22vresearch.com

Executive Summary

I have published many papers with 22V Research since our relationship began. With all the fears about a bubble weighing on investors, I wanted to provide research on the next phase of AI, when it officially enters the ROI phase. The AI investment narrative stands at a critical juncture. With over \$1 trillion deployed into data center infrastructure and mounting concerns about utilization and returns, it makes sense for investors to question whether we're in a bubble.

What makes AI different from historical bubbles is both the speed and the dependency: the products require the buildout. There can be no revenues without the massive data center infrastructure. The data centers are the brain, or in football terms, the playbook for the quarterback. But the revenues for AI, like wins for a football team, come from the quarterback's real-time decisions on the field. This is where Neural Processing Units (NPUs) become critical.

NPUs are the edge computing architecture that transforms cloud AI from expensive, centralized intelligence into distributed, real-time execution. Without NPUs, the massive GPU buildout risks becoming stranded infrastructure, brilliant models with no economical way to serve billions of users. With NPUs, the economics invert completely. By moving inference to the edge, into phones, PCs, cars, AI agents, and humanoid robots, NPUs collapse latency to milliseconds, slash serving costs by 90%, and generate the continuous utilization that makes trillion-dollar cloud investments profitable. Together, this shift marks the handoff from the first stage of AI, the cloud-driven training era that built the foundation of intelligence, to the edge era, where that hard-won knowledge is deployed everywhere, decisions happen instantly, and the trillion-dollar AI buildout finally earns its return.

This isn't about choosing between cloud and edge; it's about understanding that the cloud buildout was necessary preparation for the NPU era. Just as a quarterback needs thousands of hours in the film room before executing under pressure, AI needed the data center phase to train the models that NPUs now deploy at scale. The companies that control this edge intelligence layer, whether through silicon, software, or applications are positioning themselves to capture the compounding returns that justify the entire AI infrastructure cycle. This is the inflection point where AI capex converts to ROI.

The Catch That Changed Everything

January 10, 1982. NFC Championship. Six seconds left. Joe Montana scrambles right, three Cowboys bearing down, the pocket collapsing. The called play, Sprint Right Option, is dead. Montana keeps drifting, buying time, eyes scanning. Then instinct takes over. He lofts the ball high toward the back corner, where Dwight Clark rises and makes *The Catch*.

This wasn't improvisation. It was preparation become reflex. Montana had absorbed thousands of hours of film study so completely that when chaos arrived, he didn't freeze or look to the sideline, he *acted*. The playbook was no longer something he consulted; it was encoded in his nervous system.

This is exactly what's happening in AI right now.

For years, artificial intelligence has lived in the cloud, brilliant but distant, like coaches in a booth calling plays through headsets. Large language models trained in massive data centers can tell you what to do, but they can't react when the defense changes. Every decision requires a round trip: device to cloud, cloud back to device. That works in practice, but in the real economy where milliseconds determine outcomes and margins define winners, it breaks down.

Enter the Neural Processing Unit (NPU): AI's on-field quarterback. If GPUs built the playbook in the data center, NPUs bring it to life at the point of action. They're not just faster chips, they're the architecture that transforms AI from centralized intelligence into distributed instinct.

For investors, this shift represents the inflection point where trillion-dollar AI infrastructure spending converts into actual returns.

Why the Cloud Alone Can't Win the Game

The first act of AI was about scale. Data centers became vast training facilities, absorbing trillions of data points and producing the most sophisticated models ever created. This was necessary, you can't execute without a playbook. But building the playbook isn't the same as winning games.

The cloud model has three critical constraints:

1. Latency Kills Adoption

Every cloud query is like the quarterback running to the sideline for approval. Even at 200-300 milliseconds round-trip, it feels slow. AI that hesitates doesn't feel magical, it feels mechanical. The difference between 400ms and 40ms response time isn't incremental; it's the threshold between novelty and necessity.

2. Energy Economics Don't Scale

Running inference in the cloud is expensive. Training a model costs millions, but serving it to billions of users costs *billions*. Every query burns electricity, bandwidth, and dollars. For enterprises, AI remains a margin drag. For hyperscalers, it's a utilization problem—trillion-dollar GPU buildouts that can't possibly be saturated by cloud workloads alone.

3. The Edge Is Where Value Lives

Data is created at the edge, in phones, cars, factories, hospitals. Sending it to distant servers and back introduces not just latency but privacy risks, network dependencies, and coordination overhead. The cloud is brilliant at training, but terrible at real-time execution in context.

The breakthrough comes when intelligence moves to where decisions need to happen. NPUs enable devices to process, reason, and act locally collapsing latency to near-zero, slashing energy costs by orders of magnitude, and making AI feel instantaneous. Just as Montana didn't need the coach's approval to make The Catch, NPUs don't need permission from the cloud to act.

What NPUs Actually Do: Instinct by Design

An NPU isn't a smaller GPU. It's purpose-built silicon for *inference*, the real-time translation of trained knowledge into action.

Think of its architecture as the anatomy of instinct:

Compute Arrays: Dense grids of tiny processing units optimized for the matrix math AI requires, like neural pathways firing in parallel.

On-Chip Memory: Local storage that holds context, eliminating the need to fetch information from distant servers. This is short-term recall, knowing the coverage before the ball is snapped.

Power Management: Dynamic voltage and frequency tuning that makes continuous AI possible without draining batteries. Endurance over four quarters.

Specialized Compilers: Software that translates models trained in PyTorch or TensorFlow into optimized instructions the NPU can execute in microseconds.

Every component is engineered to shrink the gap between *knowing* and *doing*. NPUs don't just run AI faster, they enable *autonomous intelligence*. A phone can transcribe speech offline. A PC can summarize documents without internet. A car can recognize obstacles before a network ping completes. A robot can adjust its grip mid-motion.

Each of these moments is an audible—a micro-decision made at the line, powered by silicon designed for real-time cognition.

The Expansion: From Data Centers to Everywhere

For two years, the narrative has been singular: AI equals data center buildout. But that's only Act One. The next chapter unfolds at the edge, where NPUs transform every device category into an AI platform.

Smartphones: Billions of Intelligent Endpoints

Apple's A-series, Qualcomm's Snapdragon, MediaTek's Dimensity, every flagship phone now integrates NPUs handling vision, voice, and contextual inference. Each generation multiplies compute capacity (measured in TOPS—trillions of operations per second) while shrinking power draw.

The math is staggering: With 1.5 billion smartphones sold annually, every NPU upgrade creates exponential growth in global inference volume and proportional demand for memory, sensors, power ICs, and advanced packaging.

PCs: The Cognitive Refresh Cycle

Microsoft's Copilot+ initiative, powered by Intel Lunar Lake, AMD Ryzen AI, and Qualcomm X Elite, marks the first meaningful PC upgrade cycle since Windows 10. Here, NPUs become co-pilots, accelerating local transcription, image generation, and document analysis.

Investment implication: NPUs are becoming the third major compute engine alongside CPUs and GPUs, tripling silicon complexity and ASP (average selling price) per system. A market that had stagnated is entering a multi-year refresh driven by AI capability rather than incremental performance.

AI Agents: Autonomous Intelligence Demands the Edge

The Promise and Problem of Agents

The next frontier of AI isn't chatbots, it's agents. These are systems that don't just answer questions but *complete tasks*: booking travel by coordinating flights and hotels, managing your inbox by drafting responses and scheduling meetings, or researching a topic by visiting dozens of websites and synthesizing findings.

Companies like OpenAI, Anthropic, Google, and Microsoft are racing to deploy agent frameworks. The promise is profound: AI that works *for you* rather than *with you* operating autonomously in the background, handling complexity you'd rather delegate.

But there's a fundamental architectural problem: Today's agents run almost entirely in the cloud. Every decision, every step in a multi-step plan, requires a round trip to remote servers. For simple tasks this works, but for complex workflows requiring dozens of sequential or parallel decisions, latency compounds catastrophically.

Consider a travel booking agent:

1. Check calendar availability (150ms)
2. Search flights (200ms)

3. Compare prices across airlines (250ms)
4. Check hotel availability near meetings (200ms)
5. Cross-reference dietary restrictions with restaurant options (180ms)
6. Generate itinerary and send for approval (120ms)

Total: 1,100ms and that's just six steps. Real agent tasks often require 20-50+ decision points. At cloud latency, what should feel like magic feels like waiting for a webpage to load in 2005.

Then there's the cost problem. If each agent action burns \$0.01-0.10 in API calls, and you want your agent checking email, monitoring news, managing tasks, and handling logistics continuously... you're looking at \$50-500/month in inference costs *per user*. At a billion users, that's \$50-500B annually economics that make agent-as-a-service unprofitable at any reasonable subscription price.

Enter NPUs as the Agent Architecture

NPUs fundamentally change the economics and feasibility of AI agents by enabling **hybrid agent architectures**:

Tier 1 – Local Agent Orchestration (NPU):

The agent's "executive function" runs on-device:

- Task planning and sequencing
- Context maintenance (remembering what you're working on)
- Privacy-sensitive operations (reading your email, calendar, files)
- Real-time decision-making (should I interrupt the user now?)
- Coordination between apps and data sources

Tier 2 – Cloud for Heavy Lifting (GPU):

Complex reasoning and knowledge synthesis stay in the cloud:

- Deep research requiring web search across dozens of sources
- Creative generation (writing, image creation, video)
- Training updates and model improvements
- Cross-user intelligence (learning from aggregate behavior)

Tier 3 – Specialized Edge Models (NPU):

Domain-specific models run locally for instant response:

- Email classification and prioritization
- Calendar optimization
- Document summarization
- Code completion and debugging

This hybrid architecture is only possible with NPUs. The on-device processor handles the constant background orchestration, the "quarterback calling plays", while the cloud handles the occasional deep research or creative task the

“film room analysis.”

The result:

- **10-100x latency improvement** for most agent actions (local decisions in milliseconds)
- **90% cost reduction** (only complex tasks hit expensive cloud APIs)
- **Privacy preservation** (sensitive data never leaves the device)
- **Continuous operation** (agents can run 24/7 without prohibitive costs)

The Agent Deployment Timeline

The agent era unfolds in stages, each dependent on NPU capability:

2024-2025: Simple Task Agents

Email summarization and drafting, calendar scheduling assistants, document Q&A and search

NPU requirement: 40-50 TOPS (current Copilot+ PC tier)

2026-2027: Multi-App Workflow Agents

Cross-platform task completion (book travel, coordinate meetings), proactive information gathering (research topics, monitor news), personal automation (bill payments, subscription management)

NPU requirement: 75-100 TOPS (next-gen laptop/phone tier)

2027-2029: Autonomous Personal Agents

Continuous background operation with minimal user input, complex multi-day task planning and execution, learning from user behavior to anticipate needs

NPU requirement: 100-150 TOPS (distributed across multiple devices)

2029+: Multi-Agent Coordination

Your agent coordinates with other users' agents (schedule group meetings), enterprise agents manage workflows across teams, specialized agents for legal, medical, financial domains

NPU requirement: 150+ TOPS plus seamless cloud integration

Without NPUs, we're stuck in 2024, agents that are slow, expensive, and operate in discrete “sessions” rather than continuously. With NPUs, agents become truly autonomous: always-on, context-aware, privacy-preserving systems that feel less like tools and more like digital teammates.

For investors, this matters because agents represent the killer app that drives NPU adoption beyond early adopters. Just as the iPhone's value wasn't the hardware but the App Store ecosystem it enabled, NPU-equipped devices aren't valuable because they're faster, they're valuable because they enable a new category of software (autonomous agents) that simply cannot exist without them.

Microsoft isn't building Copilot+ PCs for better Zoom backgrounds. Apple isn't adding Neural Engines for slightly faster Siri. They're building the substrate for agent-based computing, the architecture where your devices don't wait for instructions but anticipate, plan, and act.

This is the quarterback not just calling audibles, but running the entire offense autonomously, reading the defense, adjusting protection, changing routes, and executing, all while the coach (cloud AI) watches from the sideline, ready to help only when truly complex decisions arise.

And just like Montana in 1982, the magic happens when preparation meets autonomy when the agent has internalized enough intelligence locally to act without hesitation.

Automotive: Intelligence in Motion

Tesla's FSD Computer, NVIDIA's Drive Thor, Mobileye's EyeQ modern vehicles contain *dozens* of neural accelerators processing sensor fusion, obstacle recognition, and decision-making. Each car becomes a mobile inference hub.

This rewires the auto supply chain: Tier 1 suppliers become AI integrators. Semiconductor content per vehicle climbs from hundreds to thousands of dollars. The automotive industry becomes an AI industry.

Humanoids: The Physical Embodiment of AI's Endgame

If smartphones put NPUs in billions of pockets, and PCs put them on every desk, **humanoids put them into the physical economy itself**, transforming AI from a productivity enhancer into a labor replacement technology.

The global labor market represents \$50 trillion in annual wages in a \$100 trillion economy. For AI to capture even a fraction of that requires something software alone cannot deliver: physical embodiment. Humanoid robots, machines that can see, manipulate, and navigate like humans represent that bridge. But unlike digital AI, physical AI cannot tolerate latency. **A robot that needs 200ms to decide how to grip an object will drop it. One that requires cloud connectivity to maintain balance will fall.**

This is where NPUs become existential rather than optional. Consider the computational demands:

Each humanoid needs to process simultaneously:

- **Vision:** Multiple cameras at 60+ fps = 10-20 TOPS
- **Motion planning:** Inverse kinematics, path optimization = 5-10 TOPS
- **Tactile feedback:** Force sensors, grip adjustment = 2-5 TOPS
- **Language/instruction understanding:** Real-time NLP = 5-10 TOPS
- **Total: 30-50 TOPS minimum, distributed across multiple NPU modules**

Compare that to:

- Smartphone NPU: 15-45 TOPS (single chip)
- PC NPU: 40-100 TOPS (single chip)
- **Humanoid: 30-50 TOPS per subsystem × 4-6 subsystems = 150-300 TOPS per robot**

This makes humanoids the most NPU-dense platforms ever built by an order of magnitude.

The economics are equally compelling. If a humanoid costs \$20-30K at scale (Tesla's stated target for Optimus):

- A manufacturing facility with 1,000 robots = \$25M capex
- Replaces 1,000 workers earning \$40K/year = \$40M annual labor cost
- **Payback period: Less than one year** (even accounting for maintenance and energy)

At that ROI, deployment accelerates exponentially:

- **2025-2026:** Pilot deployments in controlled environments (thousands of units)
- **2027-2028:** Manufacturing and logistics scaling (hundreds of thousands)
- **2029-2030:** Broader commercial adoption across industries (millions)

- **2030+:** Consumer markets open (tens of millions)

Companies like Tesla (Optimus), Figure AI, and Apptronik are racing to deploy these platforms, with early commercial pilots already underway. **Tesla's investor day on November 6th is expected to include a significant showcase of Optimus capabilities, a milestone worth watching as the humanoid market transitions from concept to commercial reality.**

More importantly, humanoids close the AI flywheel. Every robot in the field generates:

- Vision data (obstacle recognition, manipulation training)
- Motion data (balance, locomotion, efficiency optimization)
- Task success/failure data (reinforcement learning signals)

That data flows back to cloud training infrastructure, improving models that deploy to *all* robots. **This is where trillion-dollar GPU investments finally achieve utilization at scale**, not from chatbot queries, but from millions of robots learning to work.

Humanoids represent AI's crossing of the chasm from efficiency tool to labor replacement the moment when AI stops being about margin improvement and starts being about fundamental economic restructuring. And none of it works without NPUs.

The Economics: Why NPUs Unlock ROI

In AI, milliseconds are money. Power is profit. Scale is everything.

1. Latency = User Experience = Monetization

Cloud inference introduces unavoidable delay. NPUs eliminate it. When AI responds instantly, summarizing emails mid-scroll, generating images as you type, adjusting navigation as conditions change, it crosses the threshold from tool to reflex. That's when adoption accelerates and willingness-to-pay increases.

For businesses: AI features that feel instant command premium pricing. Apple Intelligence, Copilot+ subscriptions, and autonomous driving tiers all depend on low-latency inference.

2. Power = Margin Structure

Training burns megawatts. Serving billions of users burns fortunes. NPUs operate at a fraction of that power, optimized for the quantized math inference requires (8-bit, 4-bit precision instead of 32-bit training weights).

This transforms AI economics: Instead of paying per cloud transaction (OPEX), you amortize hardware cost across millions of inferences (CAPEX). Every watt saved per device compounds across billions of endpoints. For device makers, this improves gross margins. For hyperscalers, it enables services that would be unprofitable cloud-only.

3. Utilization = Infrastructure ROI

The trillion-dollar GPU buildout created massive capacity. But ROI depends on *usage*. Cloud inference alone can't saturate that pipeline, there aren't enough paying users yet.

NPUs solve the utilization problem: By pushing inference everywhere, phones, laptops, cars, robots, they generate billions of micro-inferences per second. That constant edge activity creates training data, which drives retraining workloads, which fills GPU capacity.

This is the two-sided flywheel investors need to understand:

- Edge NPUs offload inference cost and generate behavioral data
- Cloud GPUs use that data to train better models
- Better models deploy back to edge NPUs
- Rinse, repeat, compound

Without NPUs, the cloud investment risks becoming stranded infrastructure. With them, it becomes the foundation of a self-reinforcing network where every dollar of capex generates ongoing returns.

The ROI Era: When AI Becomes Profitable

The training era was about imagination, proving what's possible when machines can reason. It required enormous upfront investment with deferred returns. Investors funded a future that hadn't yet materialized.

The inference era is about execution. This is when AI steps from the lab onto the field, from cost center to profit engine.

The Hardware Multiplier

Every device category enters an AI-driven upgrade cycle:

- **Smartphones:** Annual NPU performance doubling drives continuous replacement demand
- **PCs:** First major refresh since 2015, with NPU performance as the new buying criterion
- **Automotive:** \$1,000+ of AI silicon content per vehicle, growing to \$2,000+
- **Humanoids:** Multi-chip compute platforms with margins exceeding smartphones and deployment economics that guarantee adoption

Each layer feeds the semiconductor supply chain: advanced nodes, HBM, power management, packaging, and analog components all see accelerating demand.

The Business Model Multiplier

Big Tech monetizes NPUs across multiple vectors:

Microsoft: Copilot+ subscriptions bundling AI features, higher Office 365 ASPs

Apple: Premium pricing for "Apple Intelligence" devices, expanding services revenue

NVIDIA: Ecosystem leverage, edge deployments feed cloud training demand for its GPUs

Tesla/Figure AI/Apple/Amazon: Robot fleets generating both labor productivity and training data

Google/Amazon: Reduced cloud serving costs enabling margin expansion on AI services

This is the inflection point: AI stops being a line item expense and becomes a compounding revenue generator. Every NPU installed creates ongoing value—through subscriptions, premium pricing, cost savings, or data accumulation.

Intelligence as an Asset Class

Montana's genius wasn't the scripted play, it was what happened when the script failed. He didn't think; he *reacted*. The playbook was internalized, ready to deploy the instant circumstances demanded it.

AI is reaching that moment.

The cloud phase was preseason: expensive infrastructure, limited monetization, deferred ROI. The NPU phase is game day: distributed intelligence, continuous inference, compounding returns.

From Capex to Cognitive Capital

In traditional infrastructure, each dollar buys static capacity, a server, a network node. In the NPU era, each dollar buys *productive capability*, a device that generates decisions, predictions, and creative output continuously.

Think of it this way:

- Each NPU is a capital good producing cognitive work
- Each inference is a yield event converting sunk costs into output
- Each interaction creates data—raw material for improving intelligence

This is how AI transitions from cost structure to compounding asset. Intelligence itself becomes monetizable infrastructure as fundamental to production as electricity, and as valuable to own as prime real estate.

The Strategic Repricing

When intelligence exists locally, it behaves like infrastructure. Companies gain not just efficiency but *autonomy* the ability to process, interpret, and decide without dependence on centralized compute.

That independence translates to:

- Higher operating margins (lower cloud costs)
- Faster decision cycles (no network latency)
- Competitive moats (proprietary edge intelligence)
- Compounding advantages (data flywheel)

For investors, the implication is structural: The companies that control edge intelligence platforms whether through silicon, software, or applications are positioning themselves to capture disproportionate value as AI shifts from laboratory curiosity to ubiquitous utility.

The Bottom Line for Investors

The NPU revolution represents the moment AI infrastructure spending converts into returns.

Three Investment Theses:

1. The Semiconductor Supply Chain Expands

NPUs multiply AI silicon demand beyond data centers into every device category. Winners include: leading-edge foundries (TSMC), memory suppliers (Micron, SK Hynix), packaging innovators (ASE, Amkor), analog/power management specialists (TI, ADI), test equipment providers (Teradyne), and lithography monopolists (ASML).

Teradyne benefits as a double-leverage play: NPU complexity drives automated test equipment demand (their core semiconductor testing business), while NPU-enabled edge AI accelerates adoption of their Universal Robots and Mobile Industrial Robots divisions. They're both supplier to and beneficiary of the NPU wave.

ASML represents the ultimate infrastructure dependency. Every leading-edge NPU from Apple's A18 to Qualcomm's Snapdragon 8 Elite requires 3nm or better process technology, which means extreme ultraviolet (EUV) lithography. ASML's monopoly on EUV (and upcoming High-NA EUV for 2nm and beyond) makes them the unbreakable bottleneck. If NPUs drive even a 20% increase in leading-edge wafer demand, that translates to hundreds of additional EUV tools worth \$50-100B+ over the next decade.

2. The Platform Winners Consolidate

Companies controlling edge AI platforms through silicon IP, operating systems, or application ecosystems gain pricing power and margin expansion. Winners: NVIDIA (cloud-to-edge ecosystem), Microsoft (software monetization), Apple (hardware ASPs), Qualcomm (mobile/auto), Tesla (embodied AI).

3. The Utilization Flywheel Accelerates

NPU solve the “trillion-dollar problem” making GPU infrastructure profitable by driving constant retraining demand. Winners: hyperscalers with integrated cloud-edge strategies (Microsoft Azure, AWS, Google Cloud) and pure-play AI infrastructure (CoreWeave, Lambda Labs).

The Risk if This Doesn’t Happen

Without NPUs scaling, AI remains cloud-constrained: too expensive, too slow, too limited. The infrastructure buildout becomes a stranded asset, capacity without customers, capability without monetization. Margins compress, capex returns disappoint, and the AI trade unwinds.

With NPUs, the opposite happens: Infrastructure gets utilized, inference scales economically, business models close, and compound effects accelerate. The trillion-dollar bet pays off.

Closing Thought

Montana didn’t make The Catch by calling the perfect play. He made it by being ready when the play broke down, when preparation met opportunity and instinct took over.

AI’s Catch moment is happening now. The playbook has been written in the data center. The quarterback is taking the field. From smartphones to humanoids, NPUs are enabling AI to act without hesitation, to process, reason, and execute at the edge where decisions matter most.

That’s when the trillion-dollar infrastructure bet becomes a trillion-dollar return.

The game is no longer about who has the best playbook. It’s about who can execute fastest when the pocket collapses. NPUs are how AI makes that throw and catches it.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.