22V Research
In Partnership with Analyst Hub Securities LLC, Member FINRA/SIPC

AI Macro Nexus Research

December 2, 2025

From Bits to Atoms: The AI Infrastructure Shift From LLMs to VLAs

Jordi Visser
JVisser@22VResearch.com
22vresearch.com

I listened to a long-form interview this weekend with Elon Musk. I've watched most of his interviews over the years, but this one stood out. It wasn't about quarterly results, Tesla margins or rocket design. It was Musk speaking broadly about consciousness, physics, society, money and the future of intelligence.

But in the middle of that conversation he said something that immediately caught my attention because it hit a key point from my weekly video.

"AI is going to move from thinking to acting. Real-world AI robots, autonomous vehicles, machines moving through space, this is the true frontier."

He said it casually, the way someone might comment on the weather. But it's one of the most important statements for investors thinking about AI alpha over the next three years. Because what he's really saying is: the next phase of AI is not cognitive. It's kinetic.

The entire artificial intelligence investment narrative to date in the three years since the launch of ChatGPT has been about manipulating bits, processing text, generating language, and reasoning through problems that exist purely in digital space. GPUs multiply matrices. Memory stores model weights. Networks shuffle data between processors. This architecture of cognitive AI has created one of the greatest capital expenditure cycles in history, concentrating wealth in a handful of semiconductor and cloud infrastructure companies.

To borrow a line from Good Will Hunting, this phase was AI "reading about the Sistine Chapel," not standing beneath it. In the famous park scene, Robin Williams reminds Matt Damon that knowing something intellectually is not the same as experiencing it: "You don't know what it smells like in the Sistine Chapel." That is exactly where we are with AI today. The complexity of this interaction of experiencing the physical world instead of reading about it is why I say PMIs are going to move higher. The buildout is immense to make it happen and you need to think three years from now as to where we will be rather than where we are now just like with ChatGPT in November 2022. All of this will ultimately lead to factor shifts as well within the equity market as the winners and losers change in this next phase.

This shift is happening now despite all the AI bubble talk and worries about investments and ROIC. The next phase of AI is not about thinking faster anymore, it's about moving atoms. LLMs taught AI to read. VLMs taught AI to see. VLAs will teach AI to act. Vision-Language Models (VLMs) closed the gap between text and perception, but Vision-Language-Action (VLA) models extend this into physical agency. This is not an incremental upgrade; it is an architectural fracture. Where large language models required us to build bigger datacenters, VLAs require us to rebuild the interface between computation and physical reality.

The companies that move data, perceive the environment, and execute precision motion in three-dimensional space will capture value that today's market has almost entirely overlooked. The trillion-dollar compute infrastructure was the table stakes. The multi-trillion-dollar kinetic infrastructure is the game we're about to play.

Why VLAs Break Everything

The Large Language Model era optimized for a specific problem: train transformers on massive amounts of text, then deploy them to answer questions or generate content. The workload was cloud-based and latency-tolerant, a user could wait half a second for a response. The constraint was pure compute power.

VLAs invert nearly every assumption. Instead of discrete text measured in kilobytes, VLAs ingest continuous video streams measured in terabytes. A single minute of high-resolution, multi-camera footage contains more information than millions of text tokens. Instead of generating buffered text responses, VLAs output real-time motor commands to physical actuators, robot joints, vehicle steering systems, drone rotors, where a 10-millisecond delay can mean the difference between success and catastrophic failure.

This creates what we call the physics constraint. In the LLM world, if your model hallucinates, you get a factual error, embarrassing but not catastrophic. In the VLA world, if your model hallucinates while controlling a robot, you get a “kinetic hallucination”: the AI commands a physically impossible or dangerous action. The stakes are categorically different.

VLAs introduce three new bottlenecks that GPUs alone cannot solve: the Bandwidth Wall (video ingestion overwhelms traditional interconnects), the Memory Capacity Gap (context windows explode to include environmental maps and sensor histories), and the Actuation Precision Requirement (commands must execute with sub-millisecond latency). These are not software problems. They are materials science, photonics, and mechanical engineering problems. And they create entirely new categories of infrastructure winners.

The Data Center Transformation: From Compute to Bandwidth

Training LLMs on text required massive parallelism connecting thousands of GPUs to process enormous datasets. But the data itself was relatively compact. The bottleneck was compute: how many trillions of operations per second could you sustain?

Training VLAs on video flips this dynamic. The bottleneck becomes data movement. A robotics company training a manipulation model needs to ingest continuous streams from hundreds of robots, each equipped with multiple cameras and sensors. This creates server-to-server communication within the datacenter that scales exponentially. In LLM clusters, this traffic was manageable. In VLA clusters, it becomes the primary constraint.

This is why optical interconnects are becoming the circulatory system of AI infrastructure. At the speeds required for next-generation systems, electrical signals degrade within inches due to basic physics. This forces a wholesale migration to photonics: lasers, optical fibers, and eventually co-packaged optics where the light source sits directly on the chip package.

The manufacturing complexity creates natural monopolies. Unlike electronic chips where billions of transistors are created in parallel, optical components require precise physical alignment of powered lasers to optical fibers, a serial, time-intensive process that can’t be easily automated. Companies that have mastered this operate with multi-year visibility and pricing power, because the switching cost for customers is catastrophic.

But the shift goes beyond raw bandwidth. VLAs also change the memory problem. LLMs were primarily bandwidth-constrained loading model weights as fast as possible. VLAs face a capacity problem. A robot fleet must maintain massive shared context: spatial maps, object databases, sensor histories. This context must be accessible to multiple agents simultaneously. When ten warehouse robots encounter the same obstacle, they need to query and update a shared environmental model in real-time.

This is where technologies like Compute Express Link become transformative, enabling memory pooling at rack scale. For VLA workloads with highly variable memory demand, this eliminates the need to overprovision expensive memory on every server. The companies building these memory controllers are effectively collecting a tax on the entire AI infrastructure build-out, yet most trade at modest multiples because the market still thinks of them as niche components rather than foundational infrastructure.

The Edge Revolution: Inference Moves to the Robot

LLMs could run in centralized cloud datacenters because latency was tolerant, 200 milliseconds didn't matter. VLAs cannot afford this luxury. A humanoid robot walking across uneven terrain, a drone navigating obstacles, an autonomous vehicle reacting to a pedestrian these require sub-10-millisecond decision loops. Network latency to a distant datacenter already exceeds 50-100 milliseconds, disqualifying cloud-only architectures.

This means inference must happen on the robot. Edge AI compute becomes a first-class design constraint. But edge environments are hostile: power-limited (battery life matters), thermally constrained (no datacenter cooling), safety-critical (failures cause physical harm), and cost-sensitive (you're building millions of units, not thousands of servers). The compute architectures that dominate datacenters are inappropriate for edge deployment.

Instead, we see specialized edge AI chips designed for sensor fusion, efficient inference, and real-time scheduling. These edge systems typically consume 10-50 watts rather than 500-700 watts and use conventional memory rather than expensive high-bandwidth stacks.

The memory economics shift dramatically. In the LLM era, memory upside was concentrated in exotic, expensive types stacked directly on GPUs. In the VLA era, commodity memory at the edge becomes a volume driver. If humanoid robots scale to millions of units, Tesla's stated goal, and each contains 16-64GB of memory, the demand rivals or exceeds datacenter consumption. Yet this shift is barely reflected in how investors value memory suppliers.

Better Senses: AI Needs to See the Physical World

Language models trained on text inherit all the sensory work humans already did someone wrote the text describing what they experienced. VLAs must do their own sensing, and the quality of that sensing determines success or failure.

This is why perception hardware becomes a first-order investment theme. Cameras alone, Tesla's approach for autonomous driving, force the AI to solve the depth problem computationally. A single camera provides no ground truth about distance; the model must infer depth from texture and motion. This works but is compute-intensive and brittle in edge cases like nighttime or fog.

LiDAR solves this by directly measuring distance with laser pulses, offloading a massive computational burden from the neural network. For robotics beyond automotive, warehouse robots, humanoids in homes, drones in forests LiDAR is converging on essential status. Costs have collapsed from $75,000+ a decade ago to under $500 for automotive-grade units today.

But perception extends beyond depth. Industrial robots face challenges consumer products never encounter: welding arcs that oversaturate cameras, LED lighting that causes motion artifacts, extreme reflections from metal surfaces. This drives demand for specialized industrial sensors with high dynamic range and advanced shutters. Thermal imaging provides a modality invisible to the human eye detecting overheating equipment, identifying electrical faults, locating humans in smoke. Thermal cameras, long relegated to military applications, become standard equipment in the VLA toolkit.

Event-based vision, cameras that record only pixel-level changes rather than full frames, represents a direct solution to the bandwidth bottleneck. A standard camera generates 60 full images per second whether or not anything moves. An event camera generates zero data for static pixels and microsecond-latency spikes when motion occurs, reducing bandwidth by up to 90% while cutting response time dramatically. This enables catching a ball mid-flight or dodging a swinging obstacle capabilities impossible with conventional cameras.

Muscles With Zero Slack: The Actuation Layer

This is the category that defines the atoms-versus-bits divide. Language models require no gearboxes, motors, or force sensors. Embodied AI requires all of them, at automotive-grade reliability, deployed at consumer electronics scale.

The physics of robotic motion is unforgiving. A humanoid robot with 30-40 joints needs every actuator to move synchronously both feet must step at precisely the same microsecond, or the robot loses balance. This demands deterministic networking and zero-backlash gearboxes, no "play" in the joint that accumulates positioning error. The mechanical components must execute the AI's intent with sub-millimeter precision, or the elegant intelligence is wasted on sloppy hardware.

Strain wave gears have emerged as the standard for high-precision robotic joints because they offer zero backlash with high torque in a compact form. These are not commodity components. The manufacturing process requires decades of know-how about gear geometry, material hardness, and assembly tolerances. A handful of companies globally can produce them at scale.

The cost implications are staggering. A humanoid robot with 40 joints requires 40+ precision gearboxes. At current volumes, these cost $500-2,000 per unit. This means gearbox content alone can exceed $20,000-40,000 per robot rivaling or surpassing the compute cost. In the LLM world, this category generated zero revenue. In the VLA robotics world, it becomes a multi-billion-dollar market if humanoid platforms scale to even hundreds of thousands of units annually.

But precision motion requires more than gearboxes. Robots need proprioception, the sense of where their joints are and how much force they're applying. Encoders measure angular position with extreme resolution. Force sensors at the wrist and fingertips allow VLA models to modulate grip strength, the difference between gently holding an egg and crushing it. These sensors are not optional; they're what transforms a blind, clumsy mechanical system into a dexterous agent.

The market has not yet internalized that these precision motion components are as essential to VLA scaling as GPUs were to LLM scaling. The companies manufacturing them are often small, industrial-focused, and valued as cyclical automation suppliers. Yet they sit on the critical path of the embodied AI revolution.

Why the Market Is Mispricing This Shift

The financial market's mental model of AI infrastructure is anchored in the LLM paradigm: Nvidia, TSMC, and a handful of memory suppliers captured the majority of value because compute was the bottleneck. This model worked for text-based AI, where the value chain was relatively concentrated.

VLAs shatter this concentration. The bottlenecks diversify across optical components, memory controllers, sensor physics, power, batteries and mechanical engineering. Each domain has different competitive dynamics and oligopolistic structures. The market struggles to model this because it requires understanding photonics manufacturing, automotive sensor qualification, and precision motion control supply chains, domains outside the traditional semiconductor analyst's expertise. Again, think PMIs.

There's also a timing mismatch. Datacenter optical upgrades have near-term visibility, with volume ramping in 2025-2026. Memory pooling is in early deployment. These are investable today with 12-18 month catalysts. But the actuation and perception supply chains are tied to physical robotics scaling. Humanoid robots are in prototype or low-volume production. The revenue inflection for gearbox suppliers and sensor manufacturers likely arrives in 2027-2029, not tomorrow. Markets struggle with multi-year thematic investments that lack quarterly catalysts.

Yet this is precisely where asymmetric opportunity exists. By the time humanoid robot production scales are obvious when Tesla is shipping 100,000 Optimus units quarterly, when warehouses deploy thousands of autonomous robots, the component suppliers will already be priced for growth. The time to position is when the market still thinks of these companies as niche industrial suppliers, not when they're recognized as AI infrastructure beneficiaries.

Standing Beneath the Sistine Chapel

The artificial intelligence revolution is entering its second act. The first act was about building bigger brains, more parameters, faster training, better reasoning in digital space. The second act is about connecting those brains to the physical world, moving atoms with the same facility we learned to move bits. This is what Musk meant when he said AI is moving from thinking to acting. The cognitive era taught AI to read about the Sistine Chapel. The kinetic era will teach it to stand beneath it, to reach out and touch the frescoes, to understand the world through physical interaction rather than pure description.

This is not speculative. Google's RT-2 and Physical Intelligence's π_0 already demonstrate VLA models controlling real robots using vision-language understanding. Tesla is manufacturing humanoid robots at scale, targeting tens of thousands of units in 2025 and potentially millions by 2030. Autonomous vehicle companies are deploying commercial robotaxi fleets. Warehouse automation is shifting from fixed systems to mobile, vision-guided robots. The inflection is underway; the market simply hasn't recognized it yet because it's looking at the wrong indicators.

Investors trained on the LLM supercycle are waiting for the next GPU architecture, the next incremental improvement in compute efficiency. But the next wave of value creation is orthogonal to compute. It's in the companies manufacturing optical transceivers that allow video training at scale. It's in the memory controller suppliers that enable fleet-scale context windows. It's in the LiDAR manufacturers providing ground truth for kinetic AI. It's in the gearbox companies delivering zero-backlash motion control for humanoid robots.

These companies are not tangential to the AI story, they are the AI story for the next decade. The constraint has shifted from thinking to acting, from cloud to edge, from bits to atoms. The market that recognizes this shift early, that positions in the photonic fabric, perception hardware, and actuation supply chains before they're priced as AI infrastructure, will capture asymmetric returns.

The cognitive era built the mind. The kinetic era must build the body. And the body requires an entirely different set of components, suppliers, and infrastructure providers than the mind did. The opportunity is not in making GPUs slightly better, it's in building the connective tissue, senses, and muscles that allow intelligence to manifest in the physical world.

The transition from LLMs to VLAs is not an upgrade. It is a paradigm shift. And paradigm shifts, as Musk casually observed over a weekend interview, create fortunes for those who see them coming. AI is no longer just reading about the world. It's about to touch it.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

https://brokercheck.finra.org/

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

Copyrights, Trademarks, Intellectual Property.