22V RESEARCH
In Partnership with Analyst Hub Securities LLC, Member FINRA/SIPC

AI Macro Nexus Research

January 7, 2026

The Opus 4.5 Inflection Point: When AI Crossed the Threshold

Jordi Visser
JVisser@22VResearch.com
22vresearch.com

In March 2025, Anthropic CEO Dario Amodei made a bold prediction at the Council on Foreign Relations that most observers dismissed as another sign of AI being overhyped and a bubble. Speaking about the trajectory of AI in coding and programming, he stated: "We are not far from the world. I think we'll be there in 3 to 6 months where AI is writing 90% of the code and then in 12 months we may be in a world where AI is writing essentially all of the code." At the time, we were still in the dismissive world of AI post-DeepSeek. His forecast seemed audacious, even reckless. Yet those who understand how AI frontier labs operate knew to take Amodei seriously: when he speaks publicly about future capabilities, he's not speculating, he's describing what he's already witnessing in unreleased models. The labs consistently work 3-9 months ahead of public releases which in the exponential world of AI equates to years, meaning Amodei was likely already seeing signs of the 90% threshold in internal testing when he made that March prediction.

While many people began enjoying the holidays after a long year of tariff overreaction and missing the market moves, something remarkable happened between Thanksgiving and New Year's Day that proved he hadn't been exaggerating at all, he had been conservative, perhaps even holding back the full extent of what he knew was coming.

The catalyst was Claude Opus 4.5, released by Anthropic on November 24th, and more specifically, the updated version of Claude Code that accompanied it. Roughly nine months after Amodei's CFR remarks and almost exactly at the back end of his 3-6 month window, his own company shipped a model that would validate his forecast in the most dramatic way possible. What unfolded over the holiday break wasn't just incremental progress; it was an inflection point that sent shockwaves through the AI community, captured in a cascading series of revelations from some of the industry's most respected voices. For those of you reading this who don't currently use AI in your life, your window is shutting for being relevant on many levels.

The Ten-Day Cascade

Between December 26 and January 3, the span of a single holiday week, an xAI co-founder, an Anthropic AGI researcher, a Google principal engineer, a former Google Gemini lead, and the founder of Midjourney all independently concluded that Opus 4.5 plus Claude Code had fundamentally changed what was possible. This wasn't coordinated marketing or orchestrated hype. It was organic recognition from technical leaders who suddenly found themselves experiencing capabilities they hadn't anticipated arriving for months or years.

The recognition began quietly on December 26th with a simple tweet from Igor Babushkin, co-founder of xAI and former researcher at both Google DeepMind and OpenAI. His assessment was understated: "Opus 4.5 is pretty good." But it was the response from Andrej Karpathy, OpenAI co-founder and renowned AI researcher, that signaled something more profound was occurring. Karpathy, commenting in the same week Opus 4.5 dominated AI circles replied: "It's very good. People who aren't keeping up even over the last 30 days already have a deprecated worldview on this topic."

This statement deserves emphasis. Just two months earlier, Karpathy had been talking about AGI being a decade away and coding models as performing at "elementary-grade student" levels. Now he was declaring that a single month of inattention was enough to render one's understanding obsolete. For someone of Karpathy's stature and technical depth, this represented a genuine psychological break, a recognition that the pace itself had fundamentally shifted.

The same day, Kernion, an Anthropic researcher with four years at the company, posted something even more startling: "I'm trying to figure out what to care about next. I joined Anthropic four plus years ago motivated by the dream of building AGI. I was convinced from studying philosophy of mind that we're approaching sufficient scale and that anything can be learned in a reinforcement learning environment. And so now I feel like Opus 4.5 is as much AGI as I ever hoped for and I'm not sure I know what I want to spend my waking hours focused on."

This was not a casual observer or hype merchant. This was a researcher at the company that built the model, someone who had dedicated years to achieving artificial general intelligence, now openly questioning what comes next because he believed they had essentially arrived. He elaborated: "To use Claude Code is to see Claude write arbitrary software, run into errors, reliably fix them, make helpful suggestions, and perfectly follow any given instructions."

The Creator's Testament

Perhaps the most striking validation came from inside Anthropic itself. Boris Cherny, the creator of Claude Code, provided concrete evidence that demolished any remaining skepticism about what had been achieved. Over a recent 30-day period, he reported that Claude Code using Opus 4.5 wrote 100% of his contributions to Claude Code itself: 259 pull requests, 497 commits, roughly 40,000 lines of code added and 38,000 lines removed, "with no human-written additions."

This bears repeating: the person who built Claude Code was now using Claude Code to write all the code for Claude Code. The tool had become capable of improving itself without human coding intervention. Cherny framed this as a turning point for software engineering, stating that "coding is no longer the limiting factor; the bottleneck is execution and guidance," deciding what to build, reviewing, and integrating, not typing code.

The technical details reveal the transformation's depth. Claude Code ran for minutes, hours, even days at a time across approximately 1,600 sessions and 325 million tokens, using stop hooks for long-lived tasks rather than short prompt-reply loops. This wasn't a parlor trick of generating simple scripts. This was sustained, autonomous software development at a professional level, handling the full complexity of a production codebase.

The Google Validation

If there was any doubt about whether this represented a genuine breakthrough or merely internal enthusiasm, it evaporated on January 2nd when Janna Dogen, a principal engineer at Google who leads work on the Gemini API, made a confession that stunned the industry. She tweeted: "I'm not joking, and this isn't funny. We have been trying to build distributed agent orchestrators at Google since last year. There are various options. Not everyone is aligned. I gave Claude Code a description of the problem. It generated what we built last year in an hour."

This deserves emphasis: a principal engineer at one of the world's most sophisticated technology companies, embedded in Google's own frontier AI efforts, with access to virtually unlimited resources, acknowledged that Claude Code had replicated a year's worth of her team's work in sixty minutes. When pressed for details, she clarified: "It wasn't a very detailed prompt and it contained no real details given I cannot share anything proprietary. I was building a toy version on top of some of the existing ideas to evaluate Claude Code. It was a three paragraph description."

When someone asked when Google's own Gemini would reach this capability level, her response was telling: "We are working hard right now on the models and the harness." The acknowledgment was implicit, Google was behind, scrambling to catch up to what Anthropic had achieved.

The Insider Perspective

The validation continued with Ronin Anil, a former Google engineer who had led work on the Gemini models and worked on the Google Brain team on foundational training algorithms research. He didn't mince words: "I used to be a Google engineer too leveled up all the way and feel if I had agentic coding and particularly Opus I would have saved myself the first six years of work compressed into a few months."

This wasn't speculation about future potential. This was a technical leader with deep experience in building frontier AI systems claiming that the tool would have compressed six years of his professional work into a few months. These weren't marginal productivity gains of 10-20% faster work. These were testimonials of 10-50x acceleration for complex cognitive labor.

Beyond the Technical Elite

Perhaps the most revealing commentary came from Dave Holtz, founder of Midjourney, who approached the technology not as an AI researcher but as a creator and entrepreneur. On January 3rd, he wrote: "I've done more personal coding projects over Christmas break than I have in the last 10 years. It's crazy. I can sense the limitations, but I know nothing is going to be the same anymore."

Igor Babushkin responded with a quote often attributed to Lenin: "There are decades where nothing happens, and there are weeks where decades happen." To which **Elon Musk added simply: "We have entered the singularity."**

The singularity, the hypothetical future point where technological growth driven primarily by AI becomes uncontrollable and irreversible, leading to profound and unpredictable changes in human civilization. Whether hyperbole or prophecy, the statement from Musk, who has access to frontier AI development through xAI, carried weight.

The Measurement of Progress

The subjective assessments were buttressed by objective measurement. Independent evaluators like METR and Epoch back this up: METR now pegs Opus 4.5's time horizon at around 4 hours 49 minutes, the longest of any model they've tested, while Epoch's capabilities index shows frontier progress running almost twice as fast since April 2024 as in the two years before. This metric measures the maximum duration of a task, based on human expert speed, that AI can complete successfully at least half the time. Researchers noted their current testing suite was reaching its limits for measuring these upper bounds.

The Epoch AI analysis is particularly significant: it shows that the overall rate of AI progress nearly doubled in the last two years, with a sharp inflection point in early 2024 driven largely by reasoning models and reinforcement learning. METR's data showed task horizons doubling roughly every seven months, leading to predictions that AI agents could reliably handle tasks as long as a week within the next two to four years.

Beyond Software Engineering: The Democratization of Technical Capability

The significance of this inflection point extends far beyond professional software engineers. On Christmas morning at 5 AM, after reading through many of Opus 4.5 testimonials in this paper, I decided to test the technology on a problem that had haunted me for nearly a decade. During my time at a hedge fund, I asked a team of data scientists to build a market turbulence monitoring system, a complex quantitative tool requiring sophisticated modeling and visualization. Despite their technical capabilities, the project never reached completion. The challenge wasn't just coding; it was the translation barrier between domain expertise and technical implementation. The communication cycles, iterations, and alignment required stretched timelines beyond what anyone anticipated. I did not just want a model, because the key input in the biological never ending markets is the assets to choose for model, I wanted the flexibility to control it when I wanted. We never got there.

I have zero coding background. Since incorporating my first developer into my workflow while running Morgan Stanley's Brazil office in 1997, I've always operated as the domain expert, responsible for experience-driven vision, while relying on builders to translate that vision into functional systems. The division of labor was necessary but inefficient.

That morning, I opened Claude Opus 4.5 and described what I wanted: access GitHub, find a market turbulence model that existed and was highly rated, and engage with me in Socratic mode to refine unclear aspects until we had a complete specification. I enjoyed writing this last line because it shows how much I have learned in the last year using AI daily but none of this acquired knowledge takes more than five minutes to understand. The model then generated approximately 700 lines of code and created exactly what I had envisioned. The entire process for a model took less than an hour. Later that day, I provided the specific assets I wanted to incorporate, and within three total hours of work, I had a functioning model with full visualizations, a deliverable I demonstrated on my YouTube channel that week. Most importantly, as the market's focus changes, I can tweak the assets and time period used in the model. I have a warning system based on the relationship between correlations and volatility of 100 assets.

This represents a fundamental shift in how technical capability is distributed. The bottleneck is no longer access to engineering talent or technical training. The bottleneck is domain knowledge, strategic vision, and the ability to articulate what needs to exist. For investment professionals, business leaders, researchers, and domain experts across every industry, this changes everything.

The Investment Implications

What made the Opus 4.5 moment significant wasn't just technical achievement, it was the convergence of capability with accessibility and the recognition from the industry's most knowledgeable observers that a threshold had been crossed. Amodei's March prediction about 90% of code being written by AI hadn't seemed bold because he was guessing; it seemed bold because most people couldn't imagine the trajectory he could already see in models that wouldn't be released for months.

The concrete productivity metrics are staggering: twelve months of top-tier Google engineering compressed into one hour of Claude Code scaffolding. Six years of specialized work compressed into a few months. The creator of Claude Code itself no longer writing any code manually. A non-technical investment professional building in three hours what a team of data scientists couldn't complete in years.

For investors, the implications are immediate and profound. First, the arrival of truly capable AI agents is no longer a 2027 or 2028 story, it's happening now. The timeline has collapsed. Second, the classic "build versus buy" calculation that has governed enterprise software decisions for decades has been fundamentally altered. When a domain expert can build sophisticated technical systems in hours rather than months, the economics of custom development versus off-the-shelf solutions shift dramatically.

Third, and perhaps most troubling for software incumbents, even Salesforce, one of the most AI-forward enterprise platforms, pulled back on its agent strategy in December. They blamed the capabilities but as someone who runs into this daily, AI is not a push button do it for me innovation. It is a collaborator. It is a brainstormer to help you but it needs your expertise. You lead the orchestra and it has all the instruments but needs your curiosity and guidance. School and Google search have destroyed that skill in many. When the vendors themselves blame the tool, that is a warning sign of this reality. For those trying to buy software because they are cheap and fade semiconductors, I think people are missing what these engineers have said this year but now they are saying we are here. The competitive moats around enterprise software businesses begin to look dangerously shallow in this world.

Finally, the labor market implications will dominate headlines throughout 2026 and beyond. Productivity at the corporate and GDP level will shock as labor needs decline. When six years of engineering work compresses into months, when non-technical professionals can build what previously required teams of specialists, the job displacement isn't a future concern, it's a present reality. Consumer confidence will remain pressured as workers across knowledge industries psychologically recognize their vulnerability, especially current college graduates. Voters will remain on edge as the pace of disruption outstrips society's ability to adapt. This isn't just a productivity story anymore. It's a political and social instability story that will drive policy debates, electoral outcomes, and market volatility.

This isn't about automating simple tasks. It's about fundamentally restructuring how knowledge work happens, compressing timelines that previously defined career arcs into weeks or months, and democratizing technical capabilities that were previously gated behind years of specialized training. The business world and the future of work have fundamentally changed already. Most practitioners and leaders simply aren't aware of how much. And as Karpathy warned, even thirty days of inattention is now enough to leave one's understanding obsolete. Don't wait any longer.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

https://brokercheck.finra.org/

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.