22V RESEARCH
IN PARTNERSHIP WITH ANALYST HUB SECURITIES LLC, MEMBER FINRA/SIPC

AI Macro Nexus Research

January 12, 2026

The Power Constraint: How AI's Next Phase Rewrites Infrastructure Economics

Jordi Visser
JVisser@22VResearch.com
22vresearch.com

Executive Summary

The competitive dynamics of frontier AI development have fundamentally shifted from a software race to an industrial infrastructure contest. The binding constraint is no longer algorithms or GPUs, it is the ability to deliver continuous gigawatt-scale electrical power on timelines measured in months rather than years. As Elon Musk articulated: "You've got to generate the electricity. You need transformers… you've got to convert that voltage to something the computers can digest. You've got to cool the computers." Traditional utility grid interconnection requires 36-60 months while competitive AI cycles demand operational clusters within 12-24 months. This mismatch has made "Bring Your Own Generation" (BYOG) standard practice, where hyperscalers bypass utilities by deploying dedicated onsite generation and infrastructure. When one competitor deployed 100,000 GPUs and several hundred megawatts in roughly 122 days with a roadmap toward two gigawatts it proved gigawatt-scale AI is competitively necessary. Half to two-thirds of large new data-center projects now evaluate onsite generation, with Meta's Ohio campus exemplifying the model: a 200-megawatt gas plant with $1.6 billion in dedicated infrastructure.

The BYOG transition creates a structural re-rating opportunity across energy and industrial suppliers that markets still price as cyclical businesses. The constraint decomposes into four hierarchies prime power generation, transformation equipment, thermal management, and optimization systems with transformation equipment (transformers/switchgear) the

most binding due to 100+ week lead times and oligopoly structure. Scale advantages accrue to suppliers with gigawatt capacity and balance-sheet strength, but the deeper moat lies in vertical integration delivering “molecules to megawatts” as complete solutions rather than discrete components. ExxonMobil and Chevron exemplify this advantage: production scale to allocate 400-500 MMcf/d of gas to a single campus, midstream capability to build dedicated pipelines, and balance sheets to structure 20-year contracts bundling fuel supply and generation under unified accountability. Conservative assumptions, 10-20 gigawatts of new AI capacity over five years, imply $70-100 billion in infrastructure spend. The arbitrage lies in the gap between current pricing (cyclical commodities) and emerging reality (strategic infrastructure providers with multi-decade contracts to the world’s fastest-growing capex cycle).

Investment Implications

Constraint hierarchy creates differentiated exposure. Transformation equipment (transformers, switchgear) offers the most durable pricing power through 2027 due to hardest-to-solve bottlenecks. Prime power benefits from service density moats even as capacity ramps. Thermal management has highest AI density beta but faces narrative volatility.

Vertical integration commands premium economics. Suppliers delivering end-to-end accountability capture multi-decade service annuities at utility-like margins. ExxonMobil and Chevron’s “molecules-to-megawatts” integration creates moats that component vendors cannot replicate.

Geographic advantages compound over time. Gas-rich regions (Tennessee TVA, Ohio Marcellus/Utica, Texas ERCOT) capture deployments as traditional tech hubs saturate, favoring incumbents with established infrastructure and service networks.

The re-rating is only beginning. Markets apply cyclical multiples to suppliers whose AI revenues are transitioning to multi-decade contracts with investment-grade counterparties. Companies at the intersection of scale, integration, and service density are evolving from cyclicals to infrastructure utilities.

The Binding Constraint Has Shifted

In a recent interview, Elon Musk articulated what has become the defining challenge of frontier AI development with unusual clarity. The limitation is no longer algorithmic sophistication or access to cutting-edge semiconductors. “You’ve got to generate the electricity,” Musk explained. “You need transformers… you’ve got to convert that voltage to something the computers can digest. You’ve got to cool the computers. Electricity generation and cooling are limiting factors for AI.”

He went further, describing the operational reality of training at scale: “When you do the training, the power fluctuations are gigantic. The generators want to blow up basically.” These were not theoretical concerns but lessons learned while building one of the world’s largest AI training facilities, a deployment that achieved what the industry considered impossible just months earlier.

This shift represents a fundamental phase change in how competitive advantage accumulates in artificial intelligence. The race is no longer primarily about who can design better algorithms or secure GPU allocations. It has become a contest of industrial logistics: who can deliver continuous gigawatt-scale electrical power on timelines measured in months rather than years.

The Physics of the Problem

After three years of AI progress through chip and algorithm efficiency gains, the mathematics of AI scaling have collided with the realities of electrical infrastructure in ways that create existential competitive dynamics. Modern frontier AI training clusters demand 500 megawatts to 1 gigawatt of continuous power. Competitive cycles require these facilities to be fully operational within 12-24 months. Traditional utility grid interconnection processes in major U.S. regions require 36-60 months, and in some jurisdictions approach a decade.

This mismatch is not incremental, it is categorical. If one competitor like xAI achieves operational gigawatt capacity in under two years while another waits the better part of a decade for utility connection, the laggard’s models risk permanent inferiority regardless of algorithmic sophistication or engineering talent. The competitor with power runs more training cycles, produces better models, attracts more users and revenue, and compounds that advantage into market dominance. The competitor without power simply cannot close the gap through cleverness alone.

The scale of the constraint is equally stark. A continuous two-gigawatt AI facility, the target scale for next-generation deployments, requires 400-500 million cubic feet of natural gas per day when powered by gas-fired generation. This places a single AI campus in the same energy consumption category as a large baseload power station or major petrochemical complex, with an electricity footprint comparable to 1.5-1.9 million U.S. homes.

Why "Bring Your Own Power" Became Inevitable

When industry observers watched a competitor deploy 100,000 GPUs in roughly 122 days and announce a roadmap toward one million GPUs and approximately two gigawatts of total power draw, they witnessed not just engineering execution but a strategic forcing function. The deployment proved that gigawatt-scale AI infrastructure is not just possible but competitively necessary and that the only way to achieve it on competitive timelines is to bypass traditional utility procurement entirely.

The response across the industry has been swift. Evidence suggests that half to two-thirds of large new data-center projects are now actively evaluating onsite or dedicated generation, gas turbines, reciprocating engines, fuel cells, or hybrid systems. "Bring your own power" is transitioning from edge case to mainstream design choice for facilities above 100 megawatts. Meta's expanding Ohio campus is supported by a 200-megawatt gas-fired power plant built specifically for the data center, with Williams committing roughly $1.6 billion across onsite generation and dedicated pipeline infrastructure under long-term arrangements designed to bypass traditional grid bottlenecks.

This creates a self-reinforcing dynamic that compounds rather than converges. More firm power enables larger and more frequent training runs. Better training produces superior models. Superior models attract greater user demand and revenue. Higher revenue lowers cost of capital and justifies more aggressive power procurement. More power enables even larger training runs. The cycle accelerates, and laggards cannot close the gap through incremental improvements, they must compress infrastructure timelines, which in practice means taking power delivery into their own hands.

The Infrastructure Arbitrage

When Elon Musk recently moved to acquire an overseas natural-gas power plant and physically relocate it to integrate into domestic data-center infrastructure, the transaction established a remarkable precedent. The cost of dismantling, shipping, and reassembling a gigawatt-scale plant likely exceeds hundreds of millions of dollars. Yet this expenditure becomes negligible when measured against the opportunity cost of delaying frontier model training by two to three years in a winner-take-most market.

The move demonstrates a fundamental principle: when equipment lead times from major manufacturers stretch to three to five years due to global demand, capital becomes the mechanism to compress timelines. For entities racing toward artificial general intelligence, a $500 million premium for speed represents rational allocation. Time-flexible generation capacity itself is becoming scarce, with industry reporting suggesting that flexible turbine and engine capacity, particularly from rental and fast-deployment suppliers, is increasingly being reserved by a small number of hyperscale AI developers, effectively concentrating a meaningful share of short-cycle generation capacity among a handful of customers.

The Four Vertical Constraint Hierarchies

The shift to dedicated power infrastructure decomposes into four distinct constraint hierarchies, each with different bottleneck dynamics and lead-time characteristics:

1. Prime Power Generation

At the foundation sits generation capacity, gas turbines, reciprocating engines, fuel cells, and the fuel logistics to support continuous gigawatt-scale operation. Lead times for new turbine manufacturing range from 12-24 months, and the market structure is oligopolistic, concentrating pricing power among a small number of global suppliers. The constraint extends beyond equipment to fuel logistics: pipeline permitting alone can require 18-36 months, and basis risk premiums for firm gas supply to data-center corridors are emerging as a distinct cost layer.

2. Transformation and Distribution Equipment

The most binding constraint sits not at generation but at transformation, the high-voltage transformers, switchgear, and grid interconnection equipment required to step down generation output to usable rack-level power. Lead times for large power transformers have extended beyond 100 weeks in many markets, and the oligopoly structure in manufacturing (three suppliers hold roughly 68% global share) enables 10-30% price increases in allocation markets. Electrical steel shortages and skilled labor constraints make capacity additions slow and capital-intensive, creating what multiple research reports identify as the single most difficult bottleneck to solve at gigawatt scale.

3. Thermal Management Architecture

Inside the data-center fence line, AI workload density is forcing a wholesale architectural redesign from air-based cooling to liquid cooling systems. Power densities are climbing from historical ranges of 5-10 kW per rack toward 30-80 kW per rack for AI-optimized configurations, and in some next-generation designs approach 100+ kW per rack. Air cooling becomes physically impractical at these densities, necessitating direct-to-chip liquid cooling, immersion cooling, or hybrid architectures that require complete retrofit of existing facilities and greenfield design for new ones.

This migration creates both execution risk and a distinct supply chain bottleneck. Liquid cooling systems require expertise in fluid dynamics, leak prevention, and high-density plumbing that traditional data-center operators do not possess in-house. Capacity additions for cooling distribution units, manifolds, cold plates, and heat exchangers are ramping but constrained by specialized manufacturing and commissioning timelines of 5-10 months for large deployments.

4. Optimization and Firming Layer

Once onsite generation becomes primary rather than backup power, battery energy storage systems, power conversion systems, and microgrid control software become mandatory for grid stability, peak shaving, and load management. Megawatt-scale battery deployments can compress deployment timelines to roughly 90 days, but lithium and yttrium supply chains remain a medium-term constraint, and integration with turbine/engine generation requires sophisticated controls to manage power quality, harmonics, and rapid load fluctuations, capabilities that few operators currently possess at gigawatt scale.

Green AI vs. Grey AI: The Strategic Bifurcation

The competitive dynamics are creating a strategic bifurcation in the industry. “Green AI”, slower buildouts constrained by renewable intermittency and grid planning cycles, competes against “Grey AI”, fast, gas-powered, behind-the-meter deployments that prioritize speed over sustainability commitments. For the immediate future (2026-2028), the competitive exigencies of frontier model development are forcing even sustainability-committed hyperscalers toward pragmatic acceptance of fossil fuel-based generation where speed determines market position.

Major technology companies maintain public commitments to carbon-negative or 24/7 carbon-free energy. Yet the mathematics of the AI race are pushing actual deployment decisions toward the reliability and speed that only natural gas can currently provide at gigawatt scale. Small modular reactors and other nuclear options remain on 2028-2030+ timelines for regulatory approval and first-of-a-kind construction, making them strategic options for the next generation of infrastructure but not solutions for the current competitive window.

The Re-Rating Opportunity

The investment implication is a potential fundamental re-rating of how capital markets price exposure to these constraint hierarchies. The market still largely applies old-economy cyclical multiples to the companies that supply generation equipment, electrical infrastructure, and industrial-scale thermal management as if their AI-linked revenues remain tied to GDP growth and commodity volatility rather than to multi-decade contracts with investment-grade technology counterparties.

What is emerging instead is a convergence between AI infrastructure and traditional energy and industrial infrastructure that materially changes the duration, visibility, and counterparty quality of cash flows. Under conservative assumptions, hyperscalers deploying 10-20 gigawatts of new AI-dedicated capacity over the next five years—the implied spend on generation equipment, electrical systems, and supporting infrastructure exceeds $70-100 billion, with associated gas demand measured in hundreds of millions of cubic feet per day.

The temporal arbitrage for investors lies in the gap between how these companies are categorized today and what they are becoming: strategic infrastructure providers to the highest-growth, highest-capex segment of the global economy, with revenue streams increasingly anchored by long-term service agreements and take-or-pay arrangements rather than spot commodity exposure.

Geographic Advantages and the Infrastructure Heartland

Geographic advantages are emerging based on gas infrastructure proximity and regulatory environments. Regions with robust pipeline access, Tennessee's TVA territory, Ohio's Utica and Marcellus shale basins, Texas's ERCOT market with its "connect and manage" interconnection process are capturing gigawatt-scale deployments. Traditional tech hubs in California and Northern Virginia face power saturation constraints that are pushing new frontier AI infrastructure toward the industrial heartland.

This geographic shift creates a secondary effect: it favors incumbent energy and industrial companies with established operations, service networks, and midstream relationships in these regions. The moat is not just manufacturing scale but service density. When a turbine fault occurs at 2 a.m. in a facility where every hour of downtime translates into millions of dollars of lost training cycles, having a technician on site within four hours instead of forty-eight becomes the difference between operational excellence and competitive obsolescence. At gigawatt scale, where dozens of generation units may operate in parallel, maintenance logistics, spare-parts availability, and rapid fault isolation matter as much as efficiency specifications.

Why Scale Wins in the BYOG Era

The shift toward gigawatt-scale AI campuses, each requiring hundreds of millions of cubic feet of gas per day and multi-decade reliability commitments, creates natural advantages for large integrated suppliers. A continuous two-gigawatt load pushes customers toward long-term, firm arrangements that bundle fuel supply, transportation, generation assets, and electrical infrastructure into single, highly reliable systems.

Three decisive advantages accrue to scale:

First, supply scale. Allocating 400-500 million cubic feet per day of gas to a single counterparty without destabilizing existing industrial commitments requires production at multi-Bcf/day scale, something only the largest integrated energy companies can provide while maintaining portfolio flexibility.

Second, balance-sheet capacity. Supporting multi-billion-dollar integrated projects under long-duration contracts with investment-grade technology counterparties requires financial strength that limits the supplier universe to large-cap industrials and integrated energy majors.

Third, midstream integration. The ability to invest in dedicated pipelines, compression, and processing infrastructure tailored to specific data-center clusters provides supply security that becomes existential at gigawatt scale. This "molecules-to-megawatts" integration allows conversion of volatile commodity exposure into utility-like cash flows with visibility measured in decades rather than quarters.

On the equipment side, similar advantages accrue. Large industrial manufacturers with global production footprints and dense service networks capture disproportionate value as reliability requirements intensify. The constraint extends beyond generation to electrical "balance of plant", transformers, switchgear, and high-voltage distribution systems where lead times now exceed 100 weeks and force hyperscalers to pre-order infrastructure years before site selection is finalized. Yet scale advantages in individual verticals, generation, equipment, or fuel supply tell only part of the story. The deeper moat emerges when these capabilities are integrated rather than assembled.

Vertical Integration as Execution Advantage

Scale alone does not explain why certain suppliers capture disproportionate value in the BYOG transition. The deeper competitive moat lies in vertical integration capability, the ability to engineer, deliver, and maintain complete power systems rather than selling discrete components. The complexity of gigawatt-scale AI infrastructure makes “plugging pieces together” from multiple vendors operationally prohibitive. A typical deployment requires simultaneous orchestration of fuel logistics (pipeline extensions, compression, gas processing, storage), generation equipment (turbines, emissions controls, transformers), electrical infrastructure (switchgear, power transformers, protection systems), thermal systems (cooling distribution, heat exchangers), and controls integration (microgrid controllers, energy management, power quality monitoring). Each subsystem has different vendors, engineering standards, and failure modes. Integration risk compounds exponentially as interfaces multiply.

Elon Musk, whose background in vertically integrated manufacturing at Tesla and SpaceX has shaped his approach to infrastructure, articulated this integration imperative when describing the power challenge at xAI: “You’ve got to generate the electricity. You need transformers… you’ve got to convert that voltage to something the computers can digest. You’ve got to cool the computers.” The phrase “you’ve got to” appears three times in rapid succession, not as outsourced vendor relationships but as integrated operational requirements. Musk’s instinct, honed across rockets and electric vehicles, was to treat the entire power stack as a single engineering problem requiring unified accountability. When power quality issues arise at 2 a.m. and abort a training run, hyperscalers cannot afford finger-pointing between the turbine OEM, switchgear supplier, cooling vendor, and controls integrator, they need what Musk’s organizations have always demanded: a single throat to choke.

This creates decisive advantages for suppliers who can deliver “molecules to megawatts” as complete, warranted solutions. **ExxonMobil and Chevron exemplify this integration capability at scale.** Both possess production scale sufficient to allocate 400-500 million cubic feet of gas per day to a single AI campus without destabilizing existing industrial commitments, something most independent producers cannot do. More critically, they can bundle that gas supply with midstream infrastructure investment (dedicated pipelines, compression stations, processing facilities), coordinate with generation equipment partners, and structure long-term arrangements that guarantee fuel delivery, price stability, and operational accountability under a single master agreement. This “molecules-to-megawatts” integration allows them to convert volatile commodity exposure into utility-like cash flows with multi-decade visibility, while hyperscalers gain supply security that is existential at gigawatt scale. The ability to write a single contract for 20 years of firm gas supply, pipeline capacity, and generation coordination, backed by investment-grade balance sheets and global operations experience, creates a moat that pure-play equipment vendors or independent producers cannot replicate. For hyperscalers racing to deploy gigawatt infrastructure on 18-month timelines, this integrated execution capability reduces time-to-energize by eliminating the coordination overhead that plagues multi-vendor deployments. The market still largely prices ExxonMobil and Chevron as cyclical commodity businesses exposed to oil and gas price volatility. What is emerging instead is a strategic infrastructure role: becoming the primary energy supplier to the highest-growth, highest-capex segment of the global economy, with revenue streams anchored by take-or-pay agreements with investment-grade technology counterparties rather than spot market exposure.

The Constraint Roadmap: 2026-2030

The constraint hierarchy will evolve predictably over the next four years:

2026-2027: Transformer and switchgear bottleneck intensifies. This is the hardest constraint to solve due to oligopoly structure, electrical steel supply limitations, and skilled labor scarcity. Factory expansion capex is underway but will take 24-36 months to materially increase global capacity. Allocation markets will persist, supporting pricing power.

2027-2028: Prime power capacity normalizes for modular solutions. Gas turbine and reciprocating engine suppliers are ramping capacity, and rental markets are expanding. Lead times should compress from current 18-24 months toward 12-15 months for standardized configurations, though custom large-frame turbines will remain constrained.

2028-2030: Liquid cooling supply chains mature. Thermal management bottlenecks ease as specialized suppliers scale manufacturing and commissioning expertise spreads. The retrofit wave for existing data centers creates a secondary demand pulse.

2030+: Nuclear and advanced generation enter deployment. Small modular reactors and other advanced nuclear designs begin commercial operation, offering carbon-free baseload at competitive economics. Early movers who secured regulatory approvals and site permits by 2026-2027 capture first-mover advantages in regions with stringent emissions requirements.

The Strategic Playbook That Emerges

The strategic imperatives for AI infrastructure deployment have crystallized:

Site selection is now power-first. Access to high-capacity gas pipelines, permitting feasibility, and suitable generation footprints drive location decisions ahead of fiber connectivity or tax incentives.

Generation assets are integrated into the compute stack. Power is no longer outsourced utility infrastructure but designed, financed, and operated as a core component of competitive advantage.

Fuel procurement becomes mission-critical. Gas logistics, pipeline capacity, and storage are treated as operations on par with cooling systems, networking, and accelerator procurement.

Service density trumps equipment specifications. The ability to deliver rapid fault response and maintain 99.999% uptime at gigawatt scale creates moats that persist even as equipment lead times normalize.

The conclusion is inescapable: “time to energize” has displaced “time to deploy chips” as the binding constraint on competitive advantage at the frontier of AI. The companies that solve this constraint fastest and at scale are positioning themselves not as cyclical industrials exposed to GDP growth, but as strategic infrastructure utilities tied to multi-decade contracts with the world’s fastest-growing industry.

This is the investment opportunity that Colossus made visible and that the market is only beginning to price.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

https://brokercheck.finra.org/

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.