

March 9, 2026

The OpenClaw Effect: How Agentic Software is Forcing the \$90 Trillion AI Factory Transition

Jordi Visser

JVisser@22VResearch.com

22vresearch.com

During geo-political events that lead to uncertainty and corrections in the markets, it is a great time to reassess secular trends. As was the case during the Liberation Day fears, don't lose sight of AI certainty. Regardless of how long this event disrupts markets and swings the inflation – recession pendulum back in forth in the daily commentary, AI will continue to move forward. This week, I want to use recent news and a podcast to go back to a theme I wrote about in late December that is still in the early stages.

In my December 29th paper, *The Nvidia-Groq Deal: Why AI's Next Frontier Requires Architectural Revolution Over Moore's Law*, I argued that the next phase of AI would not be decided simply by who could buy the most transistors. It would be decided by who could move data through memory, power, and latency constraints more efficiently. The Groq event was not a full acquisition. It was a non-exclusive licensing agreement, along with Jonathan Ross, Sunny Madra, and other Groq team members joining Nvidia. That distinction matters because at the time, it sent a message on AI and chips. Nvidia was not swallowing a startup to remove a competitor; it was absorbing inference-first IP and talent because inference efficiency had become strategic.

The core of the December thesis still holds. At the edge, AI runs into hard physical limits faster than the market expected. You cannot miniaturize a cloud data center into a smartphone, a car dashboard, or an industrial sensor and expect the economics to work. Power budgets are tighter, latency is visible to the user, and every trip to off-chip memory carries a real cost in watts, heat, and dollars. What has changed since December is not the physics. What has changed is where the market now sees those same constraints. The memory wall that first looked like an edge-computing problem is now showing up as the organizing principle for the next generation of AI infrastructure.

That is why Jensen Huang's keynote speech at the CES event where he highlighted the upcoming launch of the Vera Rubin platform mattered more than a normal product cycle. Rubin was presented not as a better chip, but as a rack-scale system built from six tightly coupled components, with Nvidia claiming up to a 10x reduction in inference token cost relative to Blackwell. The important point for investors is not the marketing multiple. It is the architectural admission embedded in the product: AI is no longer scaling on the back of a single processor. It is scaling through extreme co-design across compute, networking, memory, storage, and software, because the unit of competition is becoming the AI factory, not the standalone GPU.

Nvidia made the second half of that admission explicit when it introduced NVLink Fusion. For the past two years, the market has tried to frame AI infrastructure as a zero-sum contest between general-purpose GPUs and custom ASICs. NVLink Fusion says that framing is obsolete. Nvidia is now explicitly enabling hyperscalers and partners to connect custom compute to Nvidia rack-scale systems, and the first named ecosystem partners include Marvell, MediaTek, Synopsys, and Cadence. In other words, the winning architecture is no longer monolithic. It is heterogeneous. GPUs will remain critical for frontier training and fast-changing inference, but stable, high-volume workloads are increasingly likely to migrate toward custom silicon that can deliver better economics per query.

The market spent two years obsessing over who could manufacture the fastest chips, but the smartest capital is now aggressively buying the connective tissue required to link these heterogeneous components. Nvidia's recent \$2 billion direct equity investment in Coherent is a structural admission that pushing electrical signals through copper is hitting severe power and heat limits. To make rack-scale systems like Vera Rubin work, they must replace electrical circuits with

optical ones. At the facility scale, Meta's multi-year, up to \$6 billion agreement with Corning to secure optical fiber exclusively for its AI data centers highlights this same macro-level networking wall. You cannot realize the economics of the heterogeneous AI factory if your architecture is choked by the power and latency costs of copper wires and insufficient network bandwidth.

This is where the software story and the hardware story finally meet. The demand accelerant is not "AI" in the abstract. It is agentic AI. The clearest signal of this shift just came from Jerry Murdock, the notoriously press-shy co-founder of Insight Partners, a growth equity giant with over \$90 billion in AUM. He is also a member of the Board of Trustees of both The Santa Fe Institute and The Aspen Institute.

In his first-ever long-form interview on Harry Stebbings' 20VC podcast late last month (in an episode titled *The SaaS Apocalypse*), Murdock argued that the next wave of software will be driven by agents acting continuously rather than humans prompting intermittently. When a founding architect of the SaaS economy goes on the record to warn that autonomous agents will turn traditional systems of record into valueless databases, infrastructure investors need to listen whether you agree or not. That distinction matters because it changes the denominator in the infrastructure model. Traditional copilots scale with human attention. Agents scale with the number of deployed workflows and the number of actions each workflow executes. Jerry specifically mentioned the OpenClaw phenomenon and how it should be read through that lens. Its popularity, more than 100,000 GitHub stars and 2 million visitors in a single week, was not proof of monetization, but it was clear evidence of latent demand for autonomous software that can operate persistently. If you don't believe Jerry, Jensen Huang said last week, "OpenClaw is probably the single most important release of software, probably ever."

Once inference demand starts to behave like industrial throughput instead of occasional user traffic, the economics change fast. Some workloads will always need GPU flexibility. But repetitive, well-defined tasks, such as ranking, retrieval, fraud screening, code review, call center triage, enterprise search, and workflow orchestration, do not need the same cost structure as frontier model development. That is why the custom silicon story has moved from white paper to earnings release. Broadcom said last week that it now has line of sight to AI chip revenue above \$100 billion in 2027. Marvell validated this exact timeline on its March 5 earnings call. CEO Matt Murphy explicitly guided that Marvell's custom silicon business is on track to grow over 20% this year and will double next year as new hyperscaler programs hit high-volume production. Furthermore, validating the optical networking bottleneck, Murphy highlighted that their electro-optics portfolio is growing at a 50%-plus rate this year. Because of this accelerating demand for both custom compute and the connective tissue between them, Marvell lifted its corporate outlook, projecting total revenue will grow more than 30% to approach \$11 billion in fiscal 2027. Third party forecasts are moving in the same direction: Omdia sees the broader AI data center chip market reaching \$286 billion by 2030, while Counterpoint expects AI server compute ASIC shipments to triple by 2027.

The foundry and packaging layer is what keeps this from becoming just another semiconductor storytelling exercise. TSMC said advanced packaging was slightly over 10% of revenue last year and that 10% to 20% of its 2026 capital budget will go toward advanced packaging, testing, mask-making, and related capacity. That is not a side business anymore; it is becoming central to the AI stack. January 2026 revenue was up 36.8% year over year, and TSMC's February 2026 monthly sales report is due on March 10, 2026. At the same time, the memory side of the system remains tight. Reuters reported in December that SK Hynix had already sold out its 2026 output and that Micron expects memory markets to remain constrained beyond 2026. The message is straightforward: even if investors debate how quickly custom silicon takes share from merchant GPUs, both models are already competing for the same advanced manufacturing, packaging, and memory bottlenecks.

The design layer may be the cleanest way to express that broadening. Every new ASIC, every chiplet-based design, and every multi-die package increases design complexity before a single wafer is produced. Cadence's results already show that shift in the numbers: semiconductor IP revenue grew 40% year over year in the first quarter of 2025, helped by AI, chiplet architectures, and foundry ecosystem build-out. Synopsys is hosting its inaugural Converge conference on March 11–12, 2026, just days before Nvidia GTC on March 16–19, 2026, where the agenda itself now centers on AI factories, agentic AI, inference, and open models. The calendar matters because it creates a tight sequence of near-term events likely to reinforce the same narrative from different layers of the stack: more design starts, more heterogeneous systems, and more urgency around cost-efficient inference.

That is why I think this is actionable now rather than merely thematic. The market does not need to wait for a five-year vision to become investable. Broadcom and Marvell have already moved the earnings math. TSMC can validate demand again on March 10. Synopsys can frame the engineering bottleneck on March 11–12. Nvidia can reinforce the AI factory and agentic inference message at GTC on March 16–19. This is a live catalyst chain, not a distant secular story. And it argues for owning the enabling layer of the heterogeneous AI factory: Broadcom and Marvell for custom silicon, TSMC for the manufacturing and packaging choke point, and Cadence and Synopsys for the rising complexity tax embedded in every new design.

It is also important to be precise about what this thesis is not. It is not a call against Nvidia. If anything, Nvidia's own product and platform strategy validates it. Rubin lowers the cost of inference inside the factory. NVLink Fusion opens the factory to custom silicon. GTC 2026 is explicitly organized around factories, agents, inference, and open models. Nvidia is not defending a GPU-only world; it is positioning itself to orchestrate a broader one. For investors, that means a portfolio built solely around the old "one winner takes all" AI framing is likely too narrow for the next leg of the cycle.

The cleanest way to say it is this: the December memory-wall thesis has matured into a March 2026 heterogeneous AI factory thesis. The edge taught the market that intelligence is constrained by data movement, not just raw compute. The data center is now learning the same lesson at far larger scale. Agentic software increases the pressure because it turns inference into continuous workload rather than occasional interaction. Custom silicon, advanced packaging, and EDA are the investable consequence. The reason to act today is the timeline. Between Broadcom and Marvell's fresh guides, TSMC's March 10 sales report, Synopsys Converge on March 11–12, and Nvidia GTC on March 16–19, the next two weeks offer multiple opportunities for the market to keep repricing this idea from theme into estimates.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.