

March 30, 2026

Marvell and the Infrastructure of the Agentic Economy

Jordi Visser

JVisser@22VResearch.com

22vresearch.com

Why OpenClaw Signals a Major Acceleration for Marvell Technology

When markets are in a drawdown driven by geopolitical uncertainty, it becomes harder to focus on bottoms-up stories embedded within longer-term thematic trends. Macro uncertainty always trickles down in fears and a loss of focus on making money. Even so, March has again been an important month for AI progress and for a transition that is accelerating quickly. In periods when investors are derisking, some of the best opportunities come from companies that can still surprise to the upside, especially when those companies are aligned with a powerful thematic trend that remains in its early innings. Marvell is one of the names in my thematic ideas that fits this description for me.

The artificial intelligence buildout is crossing a threshold that changes the economics of infrastructure demand. The shift from chatbot AI to agentic AI is no longer a distant vision, it is happening now, and the implications for infrastructure providers like Marvell Technology are profound. When both Andrej Karpathy and Jensen Huang separately highlighted OpenClaw as a critical orchestration framework in recent public appearances, they signaled that the agentic era has moved from concept to operational reality. For Marvell, that transition represents a catalyst that investors still appear to be materially underestimating.

The OpenClaw signal matters because agentic AI changes the architecture of compute demand. A simple chatbot interaction is a single inference event. An agentic workflow is fundamentally different: it involves planning, calling tools, querying databases, checking results, using memory, coordinating multiple steps, and often rerunning parts of the process. That means exponentially more tokens, more compute, more networking, more storage activity, more CPU orchestration, and far greater infrastructure intensity than a standard question-and-answer interaction. It is a different economic object, and it requires a different infrastructure stack.

Marvell sits at the center of that stack. While Nvidia captures most of the spotlight around GPU training and inference, Marvell supplies the custom silicon, electro-optics, switches, and interconnect fabric that allow the entire system to function at scale. The company has methodically positioned itself across nearly every critical bottleneck in AI data centers, from custom XPUs to 1.6T optical DSPs to Teralynx switches and the progression in management commentary over the last three earnings calls points to an inflection that investors should not ignore.

The Earnings Progression: From Recovery to Acceleration

Three quarters ago, in Q2 fiscal 2026, Marvell delivered record revenue of \$2.006 billion, up 58% year over year. CEO Matt Murphy emphasized that growth was “being fueled by strong AI demand for our custom silicon and electro-optics products.” The tone was confident, but the story was still framed partly around recovery in enterprise networking and carrier infrastructure alongside AI growth.

By Q3 fiscal 2026, the narrative had sharpened. Revenue reached \$2.075 billion, up 37% year over year, while data center revenue climbed to a record \$1.1 billion, 98% year-over-year growth and 25% sequential growth. Management’s tone shifted from recovery to acceleration, with greater emphasis on “strong demand in AI and custom silicon.” The company also raised its data center total addressable market (TAM) estimate to \$94 billion by calendar 2028, a 26% increase from prior estimates, with custom accelerated compute alone projected to reach \$55.4 billion.

The most recent quarter, Q4 fiscal 2026, reported while eyes were on the Strait of Hormuz, crystallized the trajectory. Revenue rose to a new record of \$2.219 billion, up 22% year over year, driven by robust AI demand. Murphy's commentary became more assertive: "We anticipate that year-over-year revenue growth will quicken each quarter in fiscal 2027, propelled by ongoing strength in our data center operations, with bookings continuing to increase at a **remarkable pace**". More importantly, Marvell projected fiscal 2027 revenue growth of more than 30%, approaching \$11 billion, and fiscal 2028 revenue growth of nearly 40%, approaching \$15 billion, well above analyst consensus.

This is not just growth; it is compounding acceleration. The progression in management confidence, TAM expansion, and forward guidance suggests a business entering a multi-year growth phase driven by structural demand rather than cyclical recovery driven by AI agents and the transition to inference.

The Custom Silicon Advantage in an Agentic World

Marvell's custom silicon business is the most important part of the story. In just a few years, the company has scaled from essentially zero to around \$1.5 billion in fiscal 2026 AI-related custom silicon revenue. Marvell is now designing custom ASIC solutions for hyperscalers including Alphabet, Amazon, and Microsoft. These are not merchant chips designed to compete head-to-head with Nvidia. They are bespoke accelerators optimized for specific workloads that hyperscalers cannot easily source elsewhere. As Jerry Murdock said recently on an Insight Partners-backed podcast, autonomous agent orchestration will accelerate open-source model adoption and ASIC chip demand, and Nvidia's CUDA must remain viable for 'the ASICs explosion that's coming.

In an agentic AI world, that custom silicon becomes even more valuable. Agentic systems require continuous loops of reasoning, tool use, memory retrieval, and task coordination. Those workflows demand not only raw FLOPS, but also optimized data movement, memory hierarchy, and CPU-GPU orchestration. Marvell's custom XPU business is designed for exactly this need, giving hyperscalers tailored architectures that balance compute, memory bandwidth, and interconnect efficiency for persistent, multi-step AI workloads.

The TAM expansion Marvell announced in mid-2025 reflects that reality. Custom XPU TAM is now expected to reach \$40.8 billion by 2028, growing at a 47% CAGR, with an additional \$14.6 billion in XPU attach opportunities such as retimers, CXL controllers, and co-packaged optics. Marvell's advanced packaging platform, which enables multi-die architectures 2.8x larger than conventional single-die implementations, is already in production with a major hyperscaler. This is not distant speculation. Marvell says the platform has been qualified with a major hyperscaler and is now ramping into production.

The Interconnect and Optical Advantage

Beyond custom compute, Marvell has a strong position in the interconnect and optical layers that allow AI clusters to scale beyond a single rack. At OFC 2026 in March, Marvell showcased its 1.6T optical DSP platform, the industry's first 3nm 1.6T PAM4 optical DSP featuring 200 Gbps interfaces for AI scale-out. The company also introduced its Photonic Fabric technology platform, designed to enable multi-rack optical scale-up while meeting the reach, bandwidth, latency, and energy requirements of next-generation AI clusters. That should also position Marvell to benefit as Blackwell scales and Rubin follows.

This matters enormously for agentic AI. Unlike training workloads, which can tolerate some latency, agentic inference workloads require real-time responsiveness. An AI agent orchestrating a multi-step workflow cannot afford to wait seconds for data to move across racks. It needs low-latency, high-bandwidth fabric capable of supporting continuous read-write-execute loops. Marvell's Teralynx switches, 1.6T optical interconnects, and UEC-ready data center switches fit directly into that requirement.

Management highlighted that interconnect revenue is expected to grow by more than 50% in fiscal 2027, with switch revenue alone projected to exceed \$600 million. That growth is directly tied to the buildout of large-scale AI clusters designed for inference-heavy workloads. As OpenClaw and similar orchestration frameworks push more enterprises toward deploying agentic systems, demand for Marvell's interconnect and optical products should accelerate.

Blackwell, Rubin, and the Long Runway Ahead

Nvidia's roadmap provides additional tailwinds for Marvell. The Blackwell platform is now ramping in volume, with Nvidia reporting more than \$1 trillion in cumulative orders for Blackwell and the upcoming Vera Rubin platform through 2027. Blackwell systems rely heavily on NVLink fabric and advanced interconnects, areas where Marvell provides critical components. Nvidia's networking revenue rose 142% year over year in fiscal 2026, much of it driven by Blackwell system integration.

The Vera Rubin platform, scheduled for production in 2026 with Rubin Ultra in 2027, is designed specifically to accelerate "agentic AI, advanced reasoning and massive-scale mixture-of-experts (MoE) model inference at up to 10x lower cost per token". Rubin introduces third-generation NVLink, Inference Context Memory Storage, and rack-scale confidential computing, all of which increase the importance of best-in-class interconnect, memory controllers, and optical fabric. Marvell is well positioned to supply many of those critical enablers across the system stack.

And this is only the beginning. The AI inference market is projected to grow from \$106 billion in 2025 to \$255 billion by 2030, representing a CAGR of 19.2%. Edge inference is expected to account for 70.76% of the market in 2026, driven by demand for real-time, low-latency AI processing in IoT, automotive, and industrial applications. Marvell's portfolio is diversified and spans both cloud and edge, giving the company exposure to the full spectrum of inference growth.

Why This Matters Now

The convergence of OpenClaw's rise, Marvell's earnings acceleration, and Nvidia's Blackwell-Rubin roadmap creates a rare moment in which catalysts are aligning across the AI infrastructure stack. OpenClaw is not just another open-source project; it is a signal that the orchestration layer required for agentic systems is maturing faster than consensus appreciates. When Karpathy and Huang both spotlight the same framework in separate venues, it suggests that the agentic transition is no longer theoretical, it is operational.

Marvell's progression from 58% growth to guiding toward nearly 40% growth in fiscal 2028, combined with expanding TAMs, accelerating bookings, and deeper hyperscaler relationships, points to a company capturing structural share in the most important technology buildout of the next decade. The custom silicon business alone, essentially zero revenue just a few years ago, now \$1.5 billion and still growing, represents a rare, potentially once-in-a-cycle opportunity.

From a technical perspective, Marvell has been the best-performing stock in the SMH this month, rising 16% through Friday. What stands out is that, until now, Marvell had largely failed to participate in the broader semiconductor rally tied to the compute boom. In fact, it closed Friday at nearly the same price as its 2021 peak, meaning the stock has gone essentially nowhere for more than four years, even through the post-ChatGPT surge in semiconductors. By comparison, SMH is up more than 130% from its own 2021 peak. Technically, the setup is now improving: Marvell is trading above its 20-day, 50-day, and 200-day moving averages, with all three sloping upward. Importantly, the 200-day moving average only turned higher in December 2025, which suggests positioning is still likely lighter than in other semiconductor names.

What makes the Marvell opportunity especially attractive is the convergence of three forces: a new structural catalyst in the form of OpenClaw and Vera Rubin, a steadily improving earnings narrative that culminated in a meaningful upside surprise, and a technical backdrop that is only now turning constructive after a long period of relative underperformance. When structural change, fundamental momentum, and technical confirmation begin to align before positioning is fully rebuilt, the reward-to-risk can become especially favorable.

The market still appears to underestimate the story. Only a couple months into the rise of OpenClaw, most investors I speak with have not yet embraced the significance of it for the next phase of AI as we enter the agentic world. The breadth of Marvell's position across the AI infrastructure stack will make it a major beneficiary in this growth. The skepticism that "AI is just autocomplete" still persists, creating potential mispricing in companies like Marvell that enable the real deployment of intelligence at scale.

The agentic economy is not coming, it is already here. OpenClaw signals it. Marvell enables it. And the infrastructure spending required to support it could define the next five years of semiconductor growth.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.