

April 1, 2026

The Mythos Signal: Is Recursive Self-Improvement Here

Jordi Visser

JVisser@22VResearch.com
22vresearch.com

The AI race has seen plenty of leapfrogs over the last year, but lately it feels like Claude is leapfrogging Claude. That is what makes Mythos such an important moment. Most people had never heard of Claude Mythos until unpublished Anthropic files were accidentally exposed online, revealing a model the company would later acknowledge as its most capable yet. It entered public discussion the way major AI advances increasingly seem to: first as a rumor, then a leak, and only afterward as a confirmed reality. Before the market could frame Mythos through benchmarks, demos, or carefully managed messaging, it appeared as a glimpse of something not yet meant for broad public view. In that sense, Mythos matters not just for what it may be, but for how it surfaced: as evidence that the frontier may be moving faster in private than the public fully understands.

What makes Mythos important is not just the model itself, but the system it arrives inside. Anthropic has not been moving through a normal product cycle. It has been moving through a release sprint across models, agent workflows, developer tools, and computer use. In just a few months, the company has expanded Claude across coding, planning, enterprise integrations, security, longer context windows, and direct computer interaction. It has moved from better models to better tools to more autonomous execution. Mythos does not arrive after a pause. It arrives in the middle of a rapid-fire expansion across models, interfaces, and action layers. This is not one leap forward. It is a compounding sequence.

That is why Mythos feels different. It lands not as an isolated breakthrough, but as the next visible point on a curve that is already steepening. The market still tends to think about AI progress in model generations, as if the world waits for occasional launches and then digests them slowly. But the more relevant frame now is release cadence plus agentic usefulness plus self-correction. Anthropic's product direction increasingly emphasizes systems that can plan longer, debug better, sustain agentic tasks, review code, use tools, and operate more reliably inside larger environments. That is not just benchmark improvement. That is capability turning directly into operational output.

This is where the recursive self-improvement argument starts to feel less theoretical. The important shift is that frontier models are becoming better not only at answering questions, but at working on the very categories of tasks that improve future models and future systems: coding, debugging, tool use, computer interaction, vulnerability discovery, research assistance, and long-horizon planning. At the same time, the research world is becoming more explicit about recursive loops, whether through synthetic data pipelines, weak-to-strong generalization, multimodal agents, or systems that generate, test, and refine algorithms in a feedback cycle. None of this means a full intelligence explosion has already arrived. It does mean the architecture for compounding improvement is no longer hypothetical.

That is why Mythos has to be understood as more than a cyber story. The cyber angle is real. Reports suggest Anthropic has warned U.S. officials that Mythos could materially increase the odds of large-scale cyberattacks, and the leaked descriptions point to a model that is unusually strong at exploiting software vulnerabilities. But the deeper issue is what that implies about the underlying system. A model that can identify weaknesses, operate autonomously through technical environments, and outperform prior tools in narrow but valuable domains is moving toward a more general form of machine agency. It is not AGI in the science-fiction sense. For markets, it may be something more important: economically useful recursive pressure.

Once you see it that way, Anthropic's continuous release pattern becomes more revealing than any one benchmark. Faster releases suggest a lab that is shortening the interval between capability gains and product deployment. Better agentic performance suggests more of the workflow is shifting from human prompting to machine execution. More

computer use, code review, planning, and self-correction suggest that models are improving at the exact tasks needed to accelerate further development. In other words, what the market may be seeing as ordinary product competition could actually be the early commercial surface of recursive improvement.

The reason this matters for labor is simple: once models improve at the tasks required to improve systems, the pace of deployment into real workflows no longer has to move at a human speed.

The polite public conversation around AI still leans on augmentation. The private conversation is getting darker. Uber CEO Dara Khosrowshahi recently said that many executives privately admit the scale of disruption AI is likely to cause, but avoid saying so publicly because honesty about job displacement can scare investors and hurt fundraising. He put his own estimate starkly, saying AI may replace 70–80% of the work humans do over time, with knowledge work hit first and physical work later. Even if that estimate proves too aggressive, the significance is that a major public CEO has broken the script and acknowledged the gap between private expectation and public messaging.

The labor picture is not yet catastrophic, but it is clearly deteriorating. Over the past year, job creation has been close to zero, and the weakness looks even more pronounced in knowledge-work categories most exposed to AI disruption. Strip out healthcare, and overall job growth would already be negative. For now, the damage has shown up more in hiring than in outright layoffs. But that is exactly why the real debate is not just about how many jobs AI may eventually affect. It is about the speed of the shift. If capabilities are improving recursively and release cycles keep compressing, labor markets may lose the luxury of gradual adjustment.

This is also why Silicon Valley may not be telling us everything. It is not necessarily because executives are coordinating some grand deception. It is because the incentive structure pushes toward understatement. Public companies do not want to frighten workers, regulators, or customers. Private companies do not want to slow fundraising or invite restrictions before they scale. Frontier labs want to signal safety and optimism at the same time they are racing to deploy the most powerful systems in history. That creates a natural gap between internal belief and external language. Khosrowshahi's comments matter because they make that gap visible.

The real thesis, then, is not simply that Mythos is powerful. It is that Mythos arrives at a moment when Anthropic's release cadence, product direction, and agentic emphasis all point toward a profound compression of time. Markets are still discounting AI as a multi-year margin story, debating which software layers hollow out first and which knowledge-work sectors feel pressure by 2028 or 2029. But if capability improvement is already becoming recursive, and the interval between frontier releases is shrinking, then that timeline is fundamentally mispriced. What the market still thinks of as three years from now may be arriving much sooner than consensus expects.

That acceleration is also colliding with a macro backdrop that is already fragile. We are not entering this transition from a position of labor-market strength. Job creation has already become sluggish and increasingly narrow even before the full force of autonomous execution has arrived. Against that backdrop, the shift from human augmentation to real agentic execution should force a reset in expectations. The coming labor disruption is not a distant, gradual headwind. It is an emerging shock, and the pace of that shock may be about to accelerate materially.

If that is right, then the mistake is to wait for a ceremonial AGI finish line. Markets may be looking for a single defining announcement when the real change is already underway through compounding releases, shorter iteration loops, and economically useful agents. Mythos may not be AGI, but it is one of the clearest signs yet that the disruption curve has steepened.

And once the curve steepens, the future stops arriving in a straight line. It starts arriving all at once.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.