

June 14, 2026

Entegris Update: Vera Rubin Confirms the Missing Chemical Layer of AI Complexity

Jordi Visser

JVisser@22VResearch.com

22vresearch.com

What Has Changed Since the Original Entegris Spotlight

The original Entegris paper was built around the missing chemical layer of the AI economy. Semiconductors were getting most of the attention, but the real-world manufacturing stack required a hidden layer of high-purity chemicals, filtration, CMP, deposition materials, and contamination control.

Nothing has changed with that thought.

What has changed is that Vera Rubin is giving investors and analysts a more specific map of where the **complexity** is going. The next stage of AI infrastructure is not just about adding more accelerators. It is about designing rack-scale systems where compute, memory, networking, storage, and packaging are increasingly integrated.

That **complexity** reality is the main story because every one of those design choices increases the importance of materials science.

More HBM means more advanced memory manufacturing and packaging intensity. More memory bandwidth means tighter process tolerances. More local storage means NAND becomes more strategic to the AI system. More advanced packaging means more bonding, cleaning, CMP, underfills, substrates, and thermal-interface materials. More dense racks mean more thermal and reliability challenges.

The update is that Vera Rubin is making Entegris's opportunity more visible and the chart of the stock as it attempts all-time new highs shows investors are jumping on board.

Why Entegris Is Important to the AI Buildout

This year, the broadening out theme in semiconductors away from Nvidia has been the alpha generator. This is not speculation or a bubble, this is the reality that we have transitioned away from GPUs. Beginning with CES in January, Jensen Huang has prepared the market for the transition to tokens per watt and Vera Rubin. For this transition the bottlenecks are increasingly underneath it: memory, packaging, interconnect, thermals, substrate capacity, yield, and contamination control.

Entegris sits in those lower layers.

The company monetizes complexity, not just volume. Its opportunity grows when wafers require more process steps, tighter contamination thresholds, more filtration, more specialty chemicals, more deposition, more selective etch, more CMP, more cleaning, and more advanced packaging materials.

That makes Entegris relevant to AI because every major AI semiconductor trend increases manufacturing complexity.

Leading-edge logic requires tighter purity and more advanced materials. HBM requires more CMP, more advanced packaging, more bonding, more thermal management, and more process control. NAND layer scaling requires more etch,

deposition, cleans, filtration, and contamination control. Advanced packaging adds a new process layer that is increasingly central to AI system performance.

The old semiconductor question was: how many wafers are being started?

The better AI-era question is: how much **complexity** is embedded in each wafer?

That is where Entegris benefits. If NAND moves to higher layer counts, Entegris content per wafer rises. If DRAM moves deeper into HBM and hybrid bonding, Entegris content rises. If logic moves to more advanced nodes, Entegris content rises. If advanced packaging becomes the center of AI performance, Entegris content rises.

The company is important because AI is shifting the semiconductor value chain from simple unit growth to materials intensity. Entegris is one of the companies that helps turn AI demand into manufacturable silicon.

Why Vera Rubin Makes the Theme More Timely

Vera Rubin is the **complexity** accelerant because it shows the next phase of AI infrastructure becoming more memory-intensive, more storage-intensive, more packaging-intensive, and more thermally constrained.

This is not a small change. It is a change in the physical recipe of the AI system.

The first implication is HBM intensity. Vera Rubin pushes the industry further into HBM4 and higher memory bandwidth. HBM is not just a memory product. It is a stacked, packaging-intensive, yield-sensitive architecture. It requires more wafer capacity, more bonding, more CMP, more cleaning, more filtration, and more process control than conventional memory.

The second implication is NAND intensity. The market still tends to treat NAND as the lagging memory recovery. But Vera Rubin suggests NAND is becoming more strategic to AI systems. As GPUs require faster access to larger data sets, more local storage, and tighter coordination between compute and storage, NAND becomes less of a commodity appendage and more of a system-level requirement.

For Entegris, that matters because NAND is a content-per-wafer story. More layers mean more high-aspect-ratio etch, more deposition, more cleaning, more filtration, and more opportunities for contamination to destroy value. Even if wafer starts are not booming, higher-layer NAND can still increase Entegris content per wafer.

The third implication is advanced packaging. Vera Rubin is a rack-scale system. It is not just a chip. Performance comes from how efficiently the system moves data between GPUs, CPUs, memory, storage, and networking. That makes advanced packaging and interconnect part of the core AI performance stack.

That is exactly where Entegris matters. Hybrid bonding, underfills, substrates, thermal-interface materials, CMP, cleaning, and contamination control become more important as the package becomes more complex and more valuable.

The fourth implication is process churn. The move from Blackwell to Vera Rubin to Rubin Ultra to future architectures means the AI supply chain is entering a faster cadence of requalification, new materials, new package designs, and new yield challenges. For a toll-taker like Entegris, that churn is valuable. Every new process and new material must be qualified, controlled, cleaned, filtered, and protected.

That is the timely update. Vera Rubin makes the missing chemical layer more visible because the next AI architecture depends even more heavily on the exact bottlenecks Entegris serves.

The NAND Surprise: Storage Moves Closer to the AI System

The most underappreciated part of the Vera Rubin update may be NAND.

The consensus view is that NAND is a late-cycle recovery story. Demand improves, pricing recovers, and layer scaling resumes. That is true, but it is not the full story.

The bigger surprise is that AI is pulling NAND into the performance architecture.

As models get larger and inference workloads become more complex, the system needs faster access to larger data sets and more efficient movement between storage, memory, and compute. That means NAND may become less like cheap storage sitting outside the system and more like part of the AI memory hierarchy.

This changes the quality of NAND demand. If NAND becomes more closely tied to GPUs, high-bandwidth storage, or advanced AI system architectures, then the manufacturing bar rises. More advanced NAND, higher layer counts, and tighter packaging flows would all increase the materials and purity burden.

For Entegris, this is important because NAND layer scaling is one of the cleanest examples of content-per-wafer growth. More layers mean more etch, more deposition, more cleans, more CMP, and more filtration. A higher-layer NAND wafer can be more valuable to Entegris even without a dramatic increase in wafer starts.

In other words, Vera Rubin may turn NAND from a cyclical recovery footnote into a structural AI materials story.

HBM and Advanced Packaging: The Purity Wall

HBM is the easiest place to see the Entegris opportunity.

AI systems are increasingly memory-bandwidth constrained. HBM solves that problem, but it does so by making memory manufacturing and packaging much harder. Stacked memory, through-silicon vias, wafer thinning, hybrid bonding, and advanced packaging all create more surfaces to clean, more layers to planarize, more residues to remove, and more opportunities for defects.

This creates a purity wall.

As the value of the package rises, the cost of failure rises with it. A small contamination event can destroy an extremely expensive component. That makes filtration, purification, CMP, cleaning, and process materials less discretionary. They become yield insurance.

This is the central Entegris thesis. The company benefits when the industry cannot afford defects.

Vera Rubin raises the stakes because it increases the memory and packaging content of the AI system. The more valuable and complex the package becomes, the more valuable Entegris's role becomes.

Physical AI Extends the Duration

Vera Rubin is not just a data-center story. It also points toward the broader future of physical AI.

As AI moves from centralized cloud systems into edge devices, humanoids, autonomous vehicles, drones, industrial robots, and eventually orbital data centers, the number of intelligent machines in the world rises. Each of those machines requires some combination of compute, memory, sensors, storage, power semiconductors, connectivity, packaging, and thermal management.

This is important for Entegris because it broadens the demand base.

Edge AI may not always use the most advanced nodes, but it increases the number of chips that need to be manufactured reliably. Humanoids are essentially mobile AI systems with sensors, power electronics, memory, compute, motors, batteries, and thermal constraints. AI full self-driving vehicles are sensor-rich edge data centers on wheels. Drones and autonomous systems are distributed AI nodes operating in the physical world.

The content per wafer will vary. Some of this silicon will sit on trailing nodes, where Entegris content is lower. But the direction is clear: AI is increasing the semiconductor and materials intensity of the physical economy.

Orbital data centers are the long-tail version of the same idea. They are not a near-term revenue driver for Entegris, but they are a signal that compute demand is running into physical constraints on Earth: power, land, cooling, and permitting. Space-based compute also requires radiation hardening, hermetic packaging, novel thermal solutions, and extremely high reliability. That is a specialty-materials problem.

The near-term Entegris thesis is data centers, HBM, NAND, advanced logic, and advanced packaging. The long-term thesis is that AI moves into everything physical.

The Investment Setup

The market's mistake is likely to be treating Entegris as a cyclical semiconductor consumables name tied mostly to capex and wafer starts.

The better frame is that Entegris is becoming a structural beneficiary of AI manufacturing **complexity**.

If wafer starts rise, Entegris wins.

If content per wafer rises, Entegris wins.

If NAND adds more layers, Entegris wins.

If HBM consumes more wafer capacity and packaging intensity, Entegris wins.

If advanced packaging becomes the center of AI system performance, Entegris wins.

If AI moves into cars, robots, drones, edge devices, and orbital systems, Entegris gains a broader demand base.

This is the power of the model. Entegris does not need every AI end market to work at once. It needs intelligence to keep moving into more silicon, and for that silicon to keep getting more complex.

Vera Rubin makes that trend more visible.

Earnings Commentary

Management's own commentary reinforces why the timing of this update matters. In Q4, Dave Reeder said customers are adopting "more complex device architectures" to support "the rapid growth of AI," creating "a broad set of opportunities for Entegris," and he tied that directly to "higher content per wafer." On that same call, management said Entegris was "gaining traction in DRAM and HBM," especially in "CMP consumables and advanced packaging solutions," while also winning across "next-generation NAND" in deposition materials, CMP, and selective etch. The importance of NAND is reinforced by management's own language: across the last two earnings commentaries, Entegris mentioned NAND 39 times, underscoring that NAND is no longer just a cyclical memory recovery story but an increasingly important materials-intensity driver tied to AI storage, layer scaling, and content per wafer.

The Q1 commentary then sharpened the point: management said process complexity at advanced nodes is driving higher Entegris content per wafer, that the memory market is "structurally strong" because AI workloads are reshaping DRAM and NAND architectures, and that NAND demand is now supported not only by technology transitions but by "AI-driven storage requirements." In other words, the company is saying the same thing this paper is arguing: AI complexity is moving directly into Entegris's highest-value lanes.

Conclusion

The original Entegris thesis was that high-purity chemicals and materials science are the missing layer of the AI economy.

The Vera Rubin update is that this missing layer is becoming more important, not less.

Vera Rubin shows that the next phase of AI infrastructure is more memory-intensive, more storage-intensive, more packaging-intensive, more thermally constrained, and more dependent on manufacturing precision. That is exactly the environment where Entegris's value rises.

The cleanest investment conclusion is this:

Entegris is not a bet on semiconductor units. It is a bet on the rising materials complexity required to manufacture intelligence.

AI does not just need more chips. It needs more yield, more purity, more filtration, more CMP, more deposition, more etch, more packaging, more memory, and more advanced materials.

Vera Rubin confirms that the AI buildout is moving deeper into Entegris's lane.

DISCLOSURES AND DISCLAIMERS

Analyst Certification

The analyst, 22V Research Group, primarily responsible for the preparation of this research report attests to the following: (1) that the views and opinions rendered in this research report reflect his or her personal views about the subject companies or issuers; and (2) that no part of the research analyst's compensation was, is, or will be directly related to the specific recommendations or views in this research report.

Analyst Certifications and Independence of Research.

Each of the 22V Research analysts whose names appear on the front page of this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views of in this Report.

22V Research (the "Company") is an independent research provider. The Company is not a member of the FINRA or the SIPC and is not a registered broker dealer or investment adviser. 22V Research has no other regulated or unregulated business activities which conflict with its provision of independent research.

22V Research, LLC is a professional services and independent publication organization. 22V Research, LLC is not a securities broker-dealer, not a member of the Financial Industry Regulatory Authority (FINRA), not a registered investment advisor (RIA) and not a member of SIPC.

Securities transactions, when offered, are offered by 22V Securities, LLC through LPS Capital, LLC. Certain employees of 22V Securities, LLC are dually registered as securities representatives of LPS Capital, LLC or Analyst Hub Securities, LLC. 22V Securities, LPS Capital and Analyst Hub Securities are members FINRA, SIPC.

<https://brokercheck.finra.org/>

Current Ratings Definition.

SECTOR OUTPERFORM: An "outperform" rating anticipates the company will outperform the S&P Regional Banking Index (peer group).

SECTOR PERFORM: A "market perform" rating anticipates the company will perform in line with the S&P Regional Banking Index (peer group).

SECTOR UNDERPERFORM: An "underperform" rating anticipates the company will underperform the S&P Regional Banking Index (peer group).

Limitation Of Research And Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of 22V Research Group. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material or for any damage or lost opportunities resulting from use of this data. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision. 22V Research Group shall not be liable for any direct or indirect, incidental or consequential loss or damage (including loss of profits, revenue or goodwill) arising from any investment decisions based on information or research obtained from 22V Research Group.

Reproduction And Distribution Strictly Prohibited.

No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.

22V Research Group, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of 22V Research Group is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by 22V Research Group unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to, copyright, trademark, trade secret or patent laws.